Develop Azure compute solutions – Virtual Machines

Hosting options for a web application

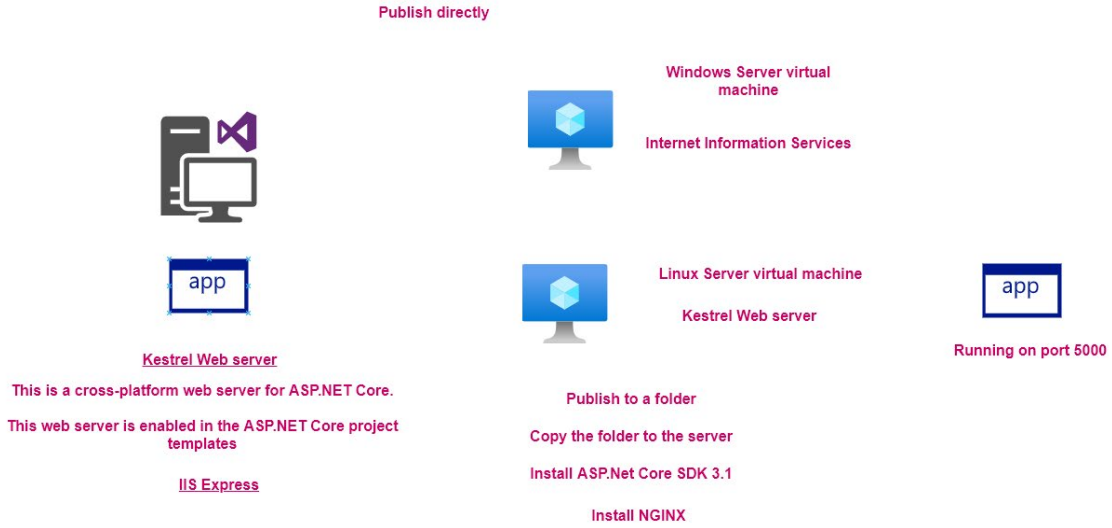

Publishing a .Net project onto a Windows server

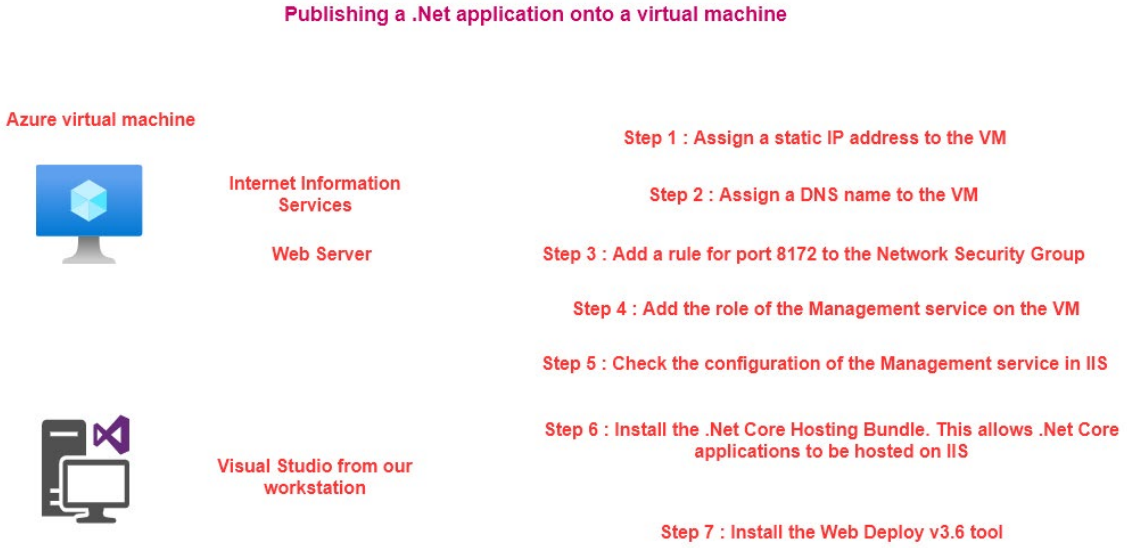

ARM templates

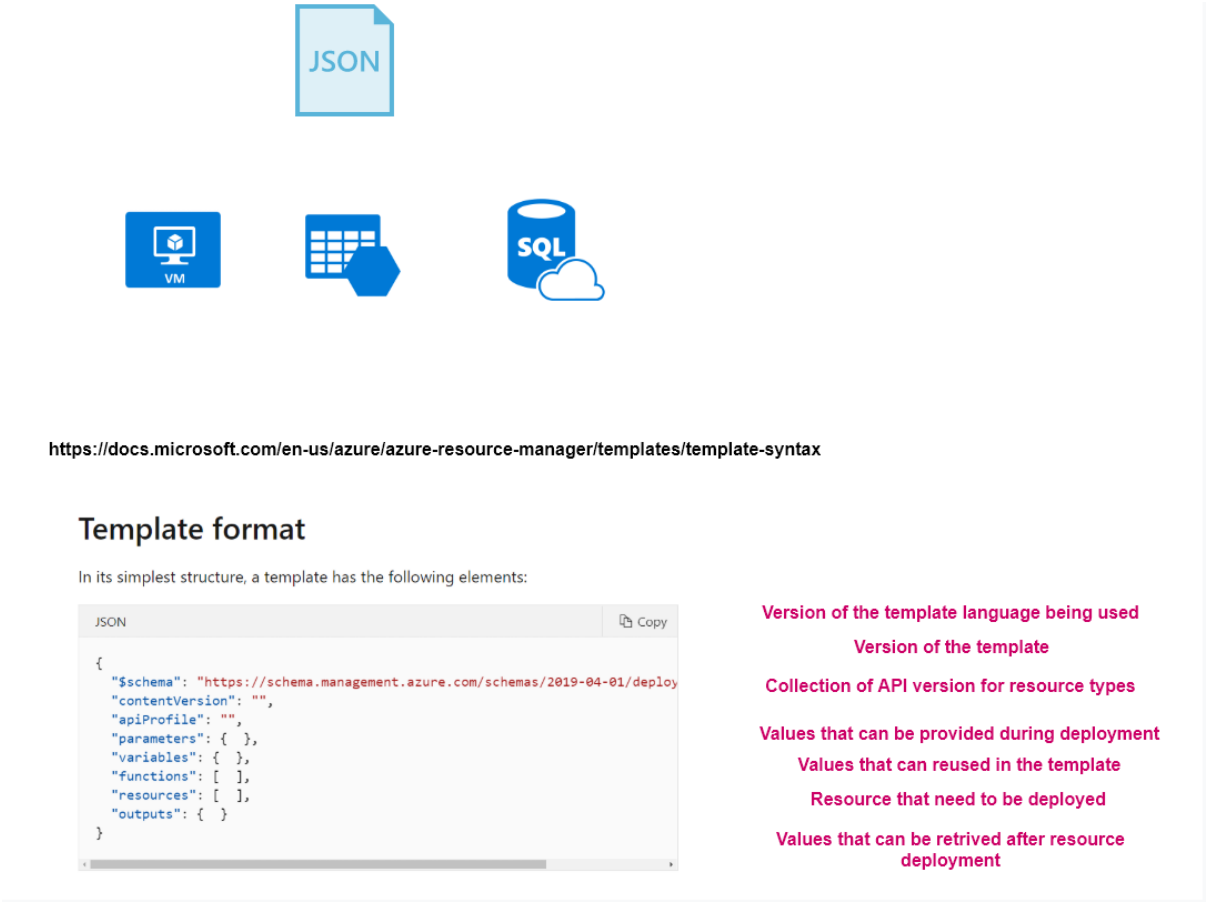

Understanding the different components of a VM deployment when it comes to ARM templates

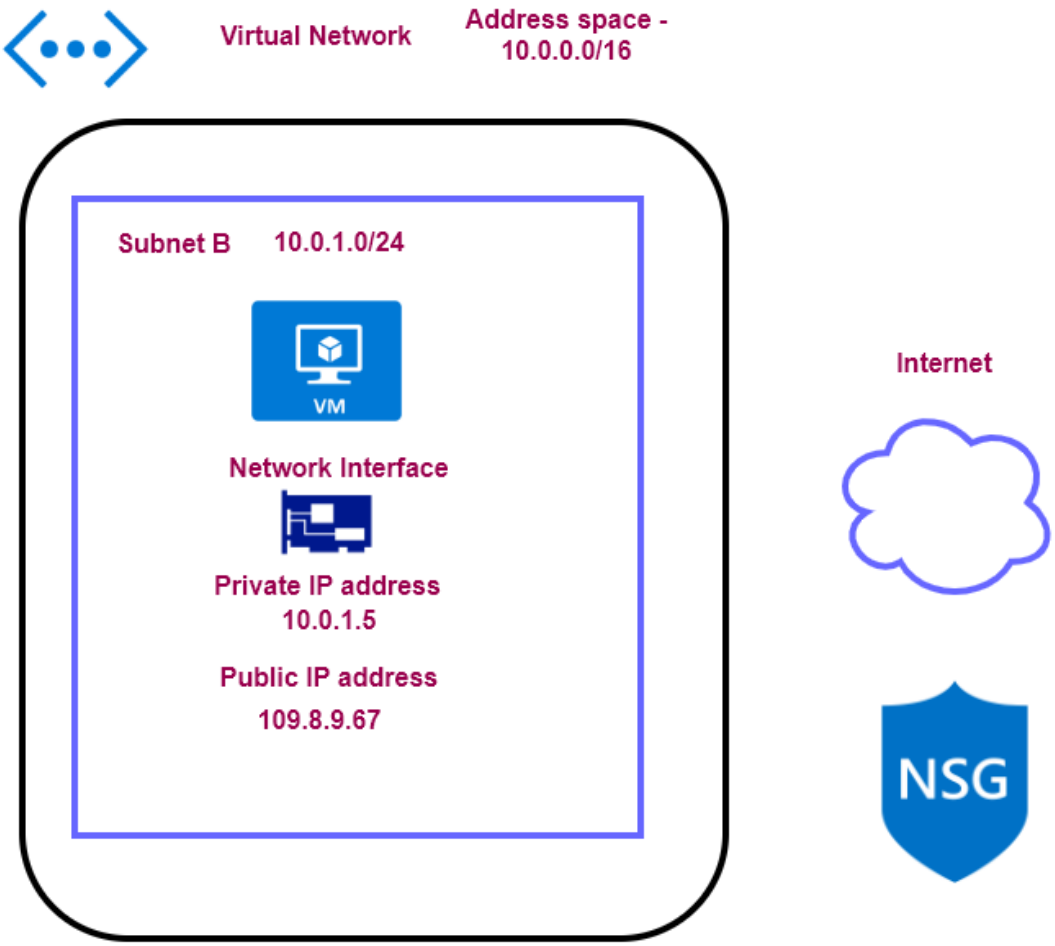

Creating an image of a VM

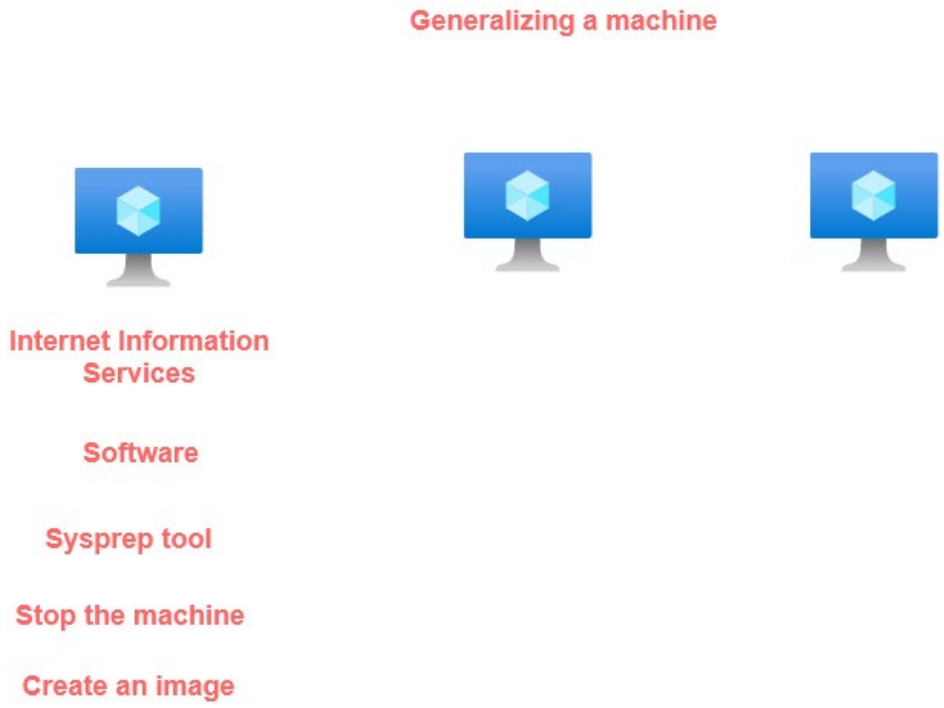

Develop Azure compute solutions – Azure Web Apps and Azure Functions

Azure App Configuration

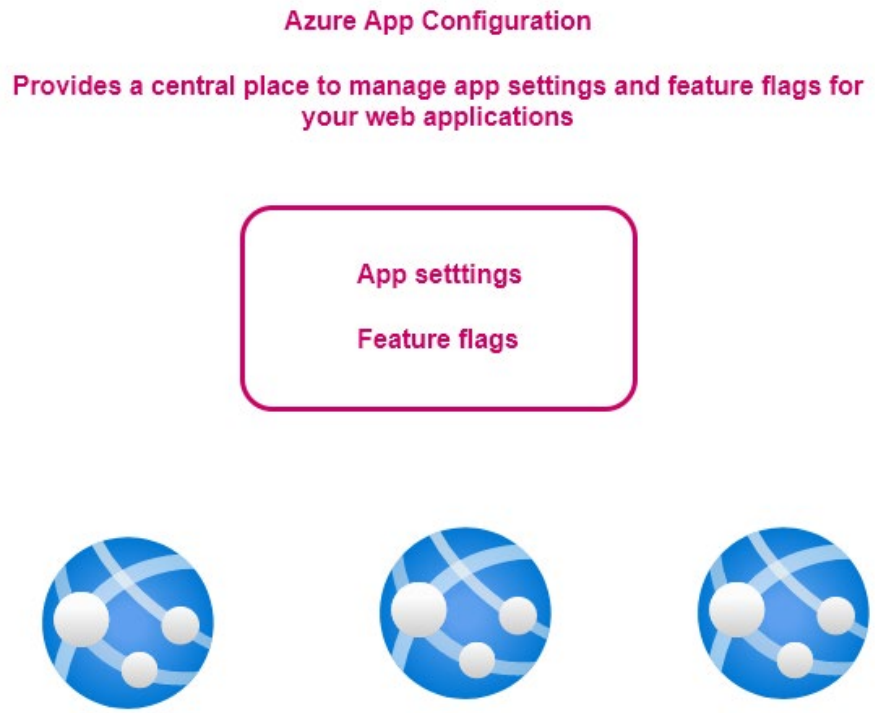

Azure Web App Logging

- You get a set of logging features that are available for the Azure Web App.
- The different types of logging that are available are
- **Application Logging** – This captures log messages that are generated by your application code.
- **Web server logging** – This records raw HTTP request data.
- **Detailed Error Messages** – This stores copies of the .htm error pages that would have been sent to the client browser.
- **Deployment logging** – These are logs when you publish content to an application.
- You can also stream logs in real time.

Durable Functions

Durable Functions

Stateless in nature

Workflow

Take customer order

Process payment

Generate receipt

Fan-out scenario

Starter Function

Orchestrator Function

Activity Function

Azure Web App – CORS

CORS - Cross-origin resource sharing

This is a method of restricting access to resources on a web page when it is requested from another domain outside the domain from which the initial resource is being served

https://courseapp1000.azurewebsites.net/api/Course

https://consumeapi1000.azurewebsites.net/demo.html

Develop Azure compute solutions – Docker, Azure Container Instances, Kubernetes

Docker Architecture

What is Docker

This is an open platform that is used for developing, shipping and running applications.

Docker has the ability to package and run an application in a loosely isolated environment called a container

Container

Isolated and lightweight environment for hosting your application

Docker host

CPU, Memory , Network

Image

This is a read-only template with instructions that are required to create the Docker container

Container

This is a runnable instance of an image

Installing Docker on a Linux VM

This manages the complete container lifecycle on the host system

containerd

Docker engine

Docker CLI

Please run a Docker container

Run your Linux systems on your windows machines

Windows subsystem for Linux

Windows system

Linux system

Windows system

Here you get access to a Linux environment running on Windows

Developers can test out their applications locally itself on the Linux subsystem.

When considering using Windows containers

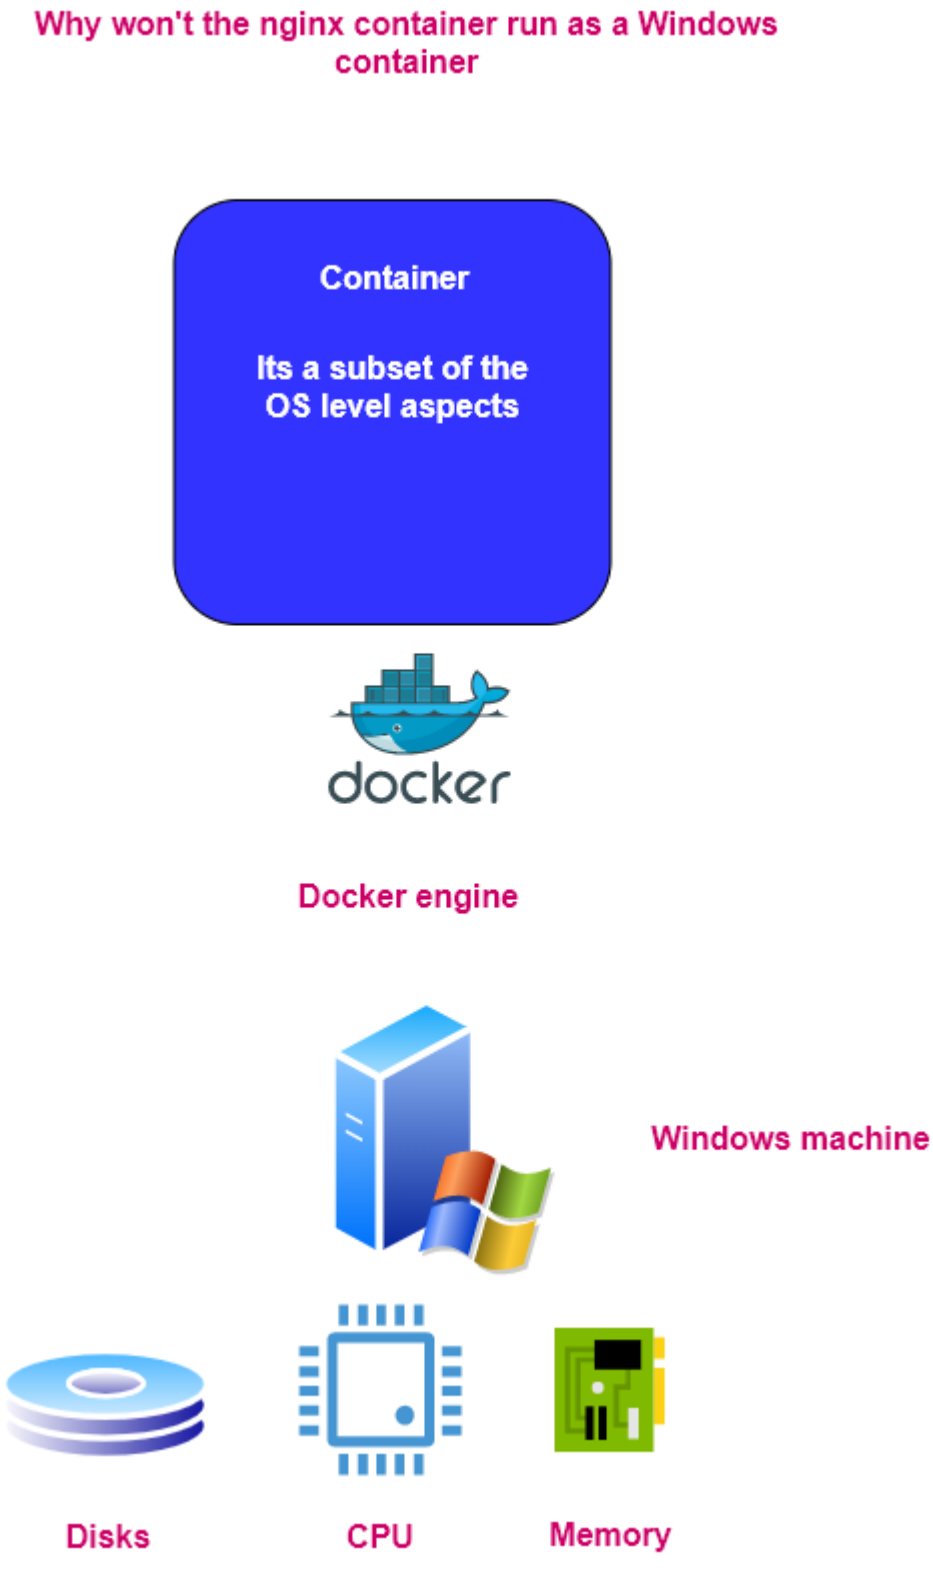

The need for Azure container registry

Repository of images

Docker engine

.Net Core application

Image

container

nginx image

nginx container

Azure Container Instances

Azure Container Instance

This service provides a fast and simple way to deploy your containers

Here you don't need to provision VM's to host your containers

Containers also get their own fully qualified domain name and IP address

Azure Container Registry

Understanding images

Repository of images

Docker engine

nginx image

nginx container

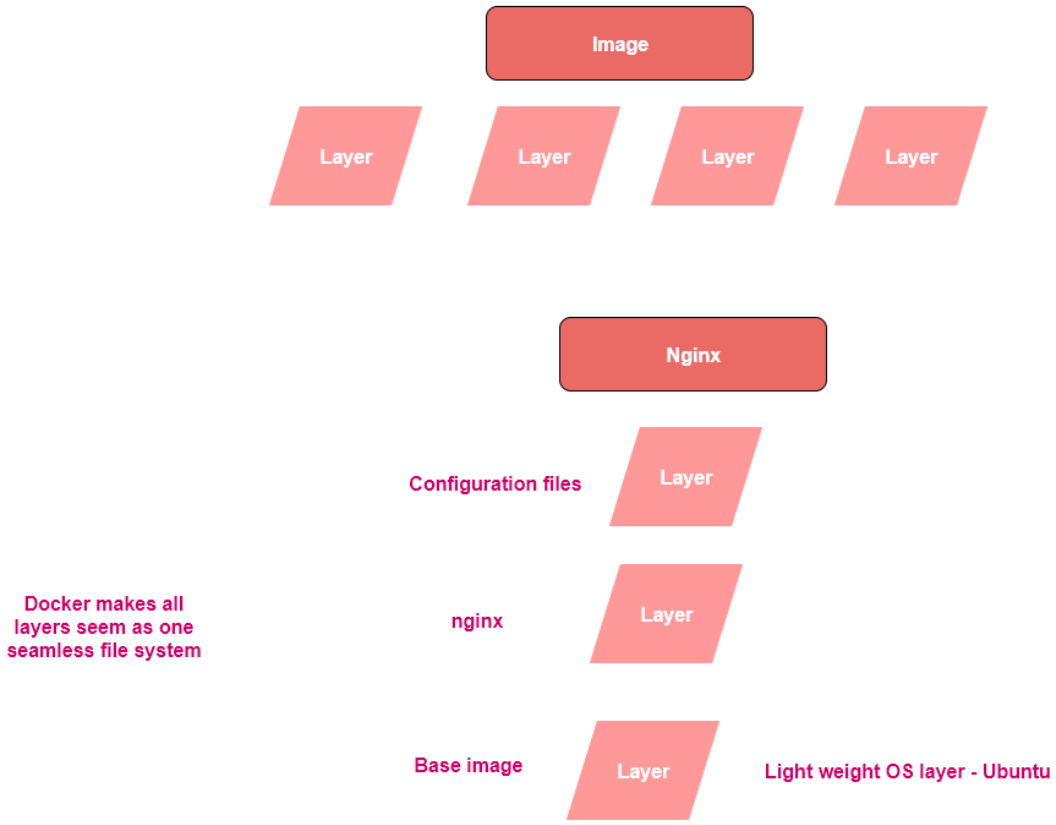
Image
Layer
Layer
Layer
Layer
Nginx
Configuration files
Layer
Docker makes all layers seem as one seamless file system
nginx
Layer
Base image
Layer
Light weight OS layer - Ubuntu

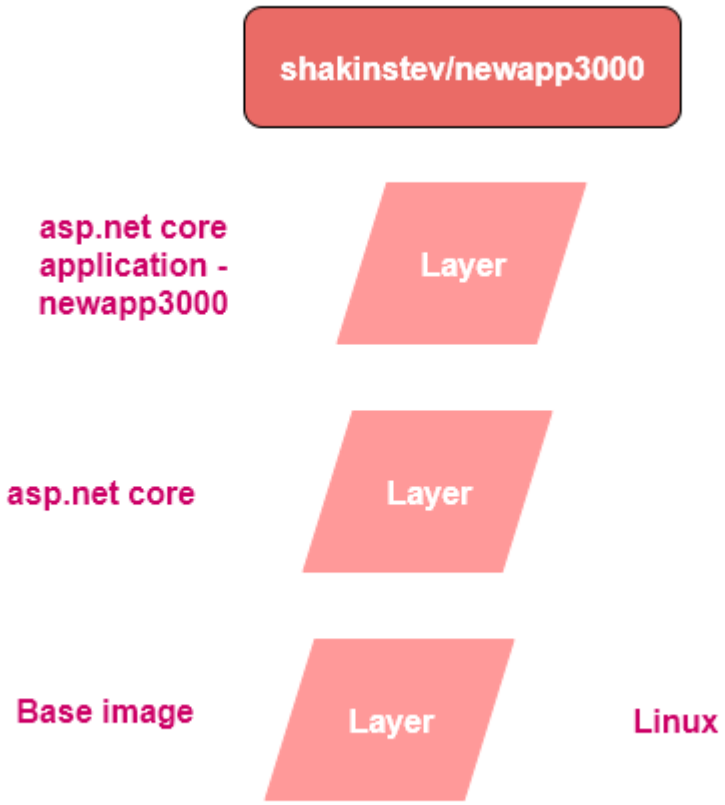

The Docker build command

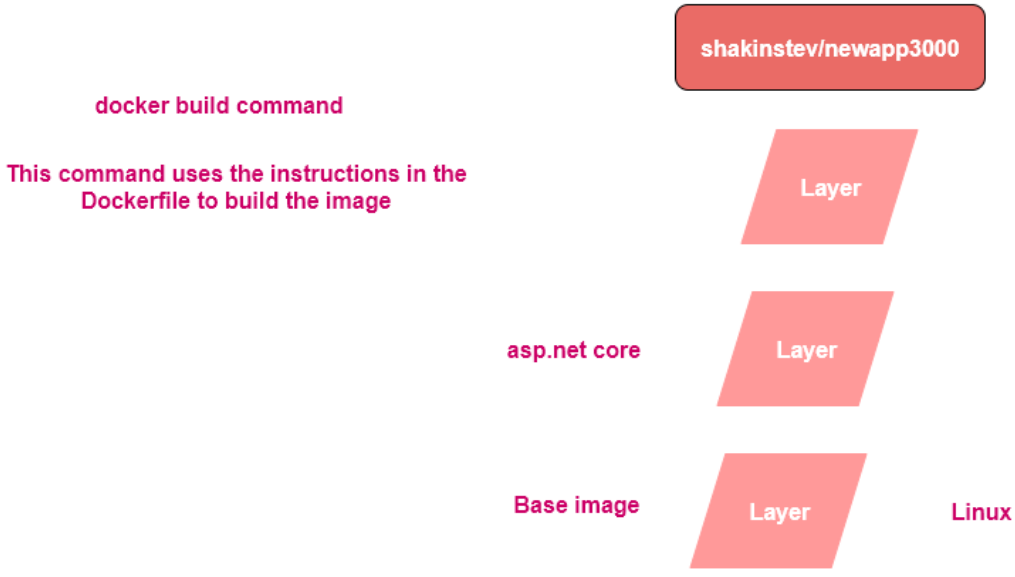

You can use volumes to persist data

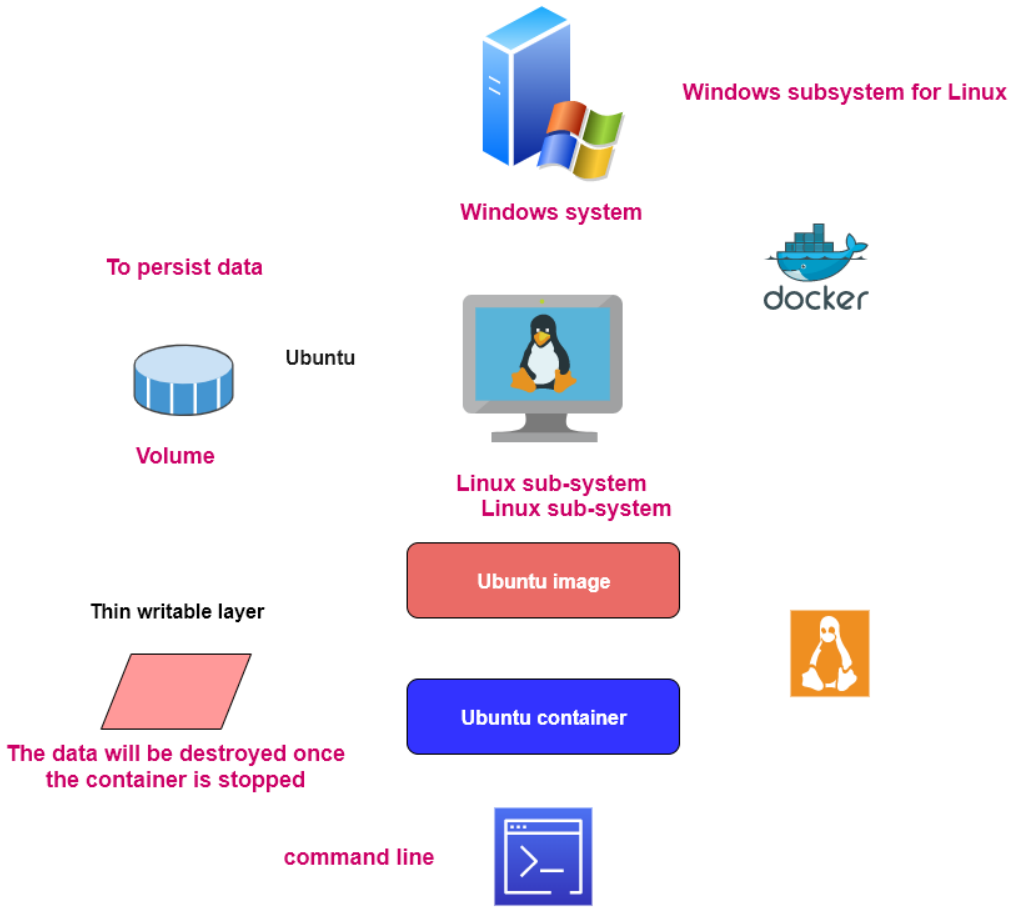

Docker compose

Docker compose

This is a tool that is used for defining and running multi-container Docker applications

You can also use a YAML file to configure the application service

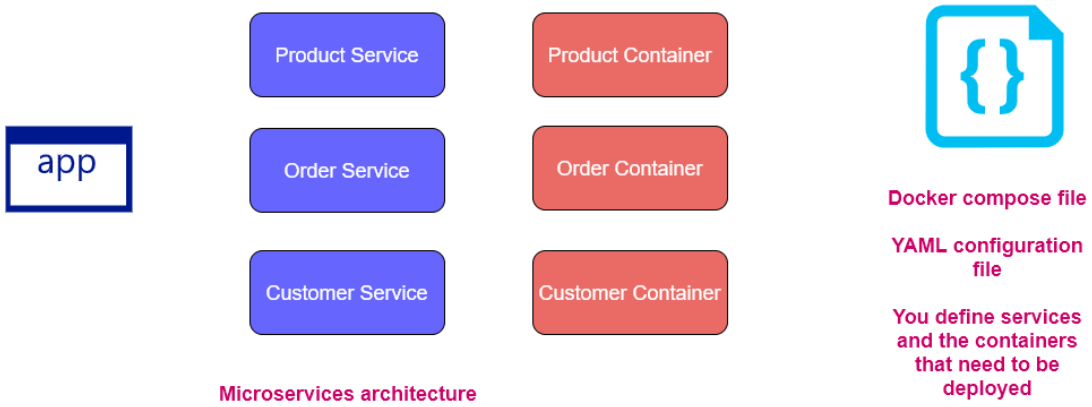

Azure Container Groups

Azure Container groups

This is a collection of containers that get scheduled on the same host machine

The containers in the container group shared the lifecycle, resources, local network and storage volumes.

The dpeloyment of the container group can be done via a Resource Manager template of a YAML file.

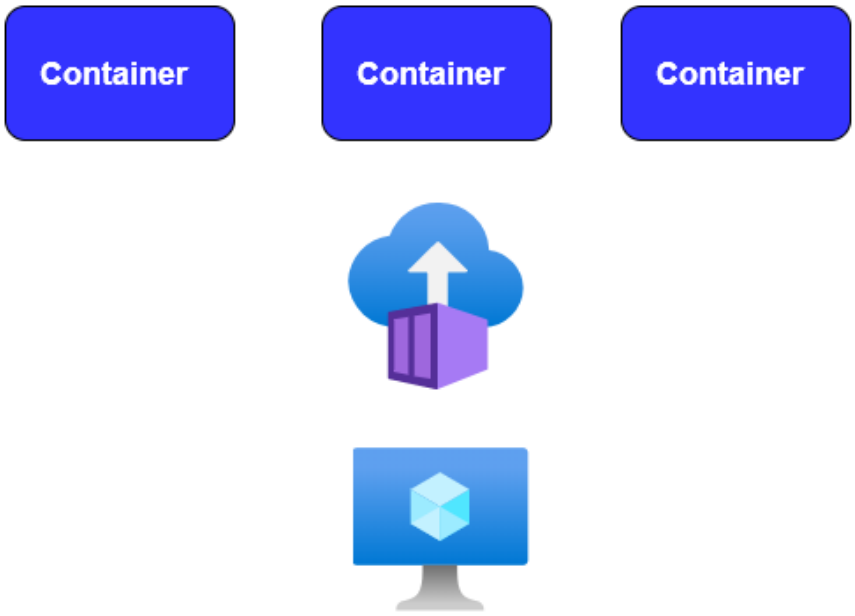

Azure Container Instances – Connecting to file shares

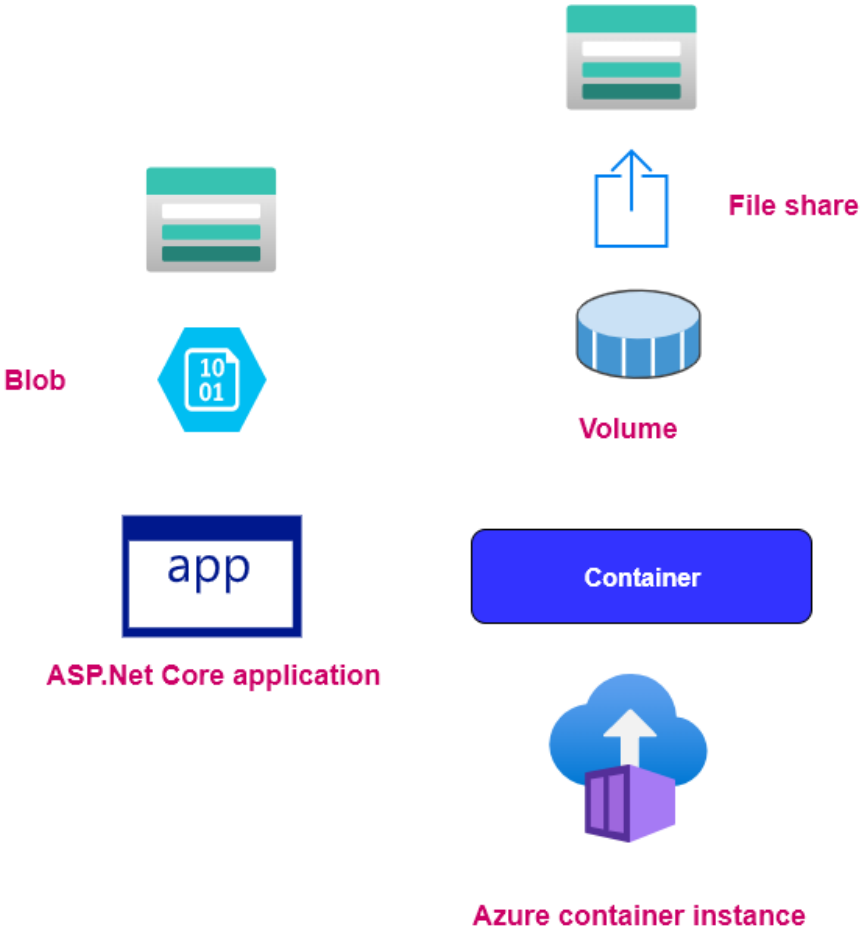
File share
Blob
10
01
Volume
app
Container
ASP.Net Core application
Azure container instance

Kubernetes

Kubernetes

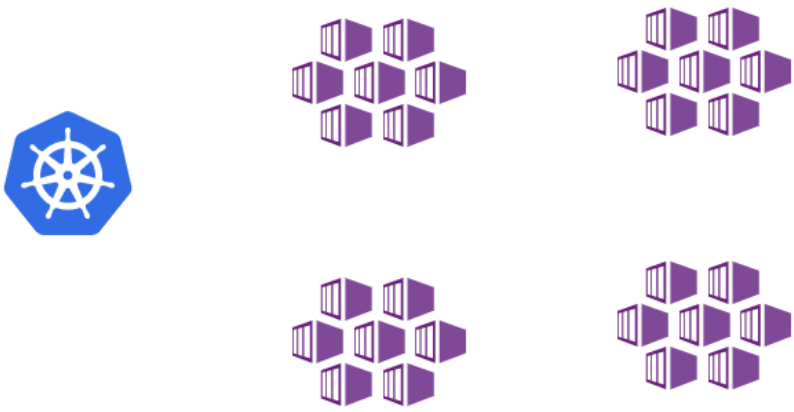

Managing containers at scale

Azure Kubernetes - Managed service for Kubernetes on Azure

Kubernetes is used to orchestrate your containers for hosting your applications

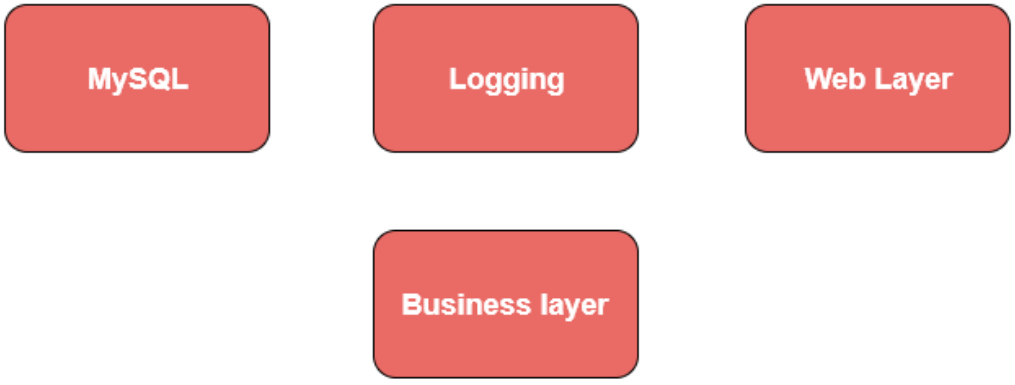

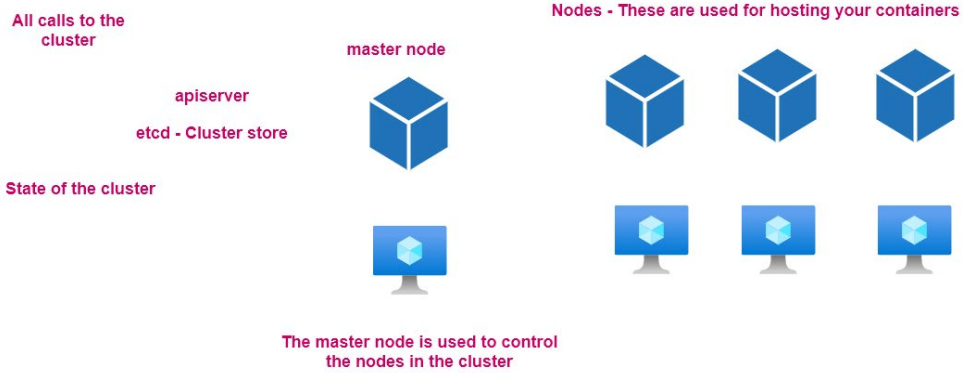

Node

kubelet - This is a kubernetes agent that runs on the node.

kubelet - It registers the node with the master node

kubelet - Will take commands from the master node for the deployment of containers

Container runtime - This is used to actually taking the images and deploying the containers on the node

Kube-proxy is used for managing the networking aspects for the containers

Azure Kubernetes Pods

Deployment of containers

Pod

A Pod is used to group one or more containers.

The pod gets shared storage and network resources

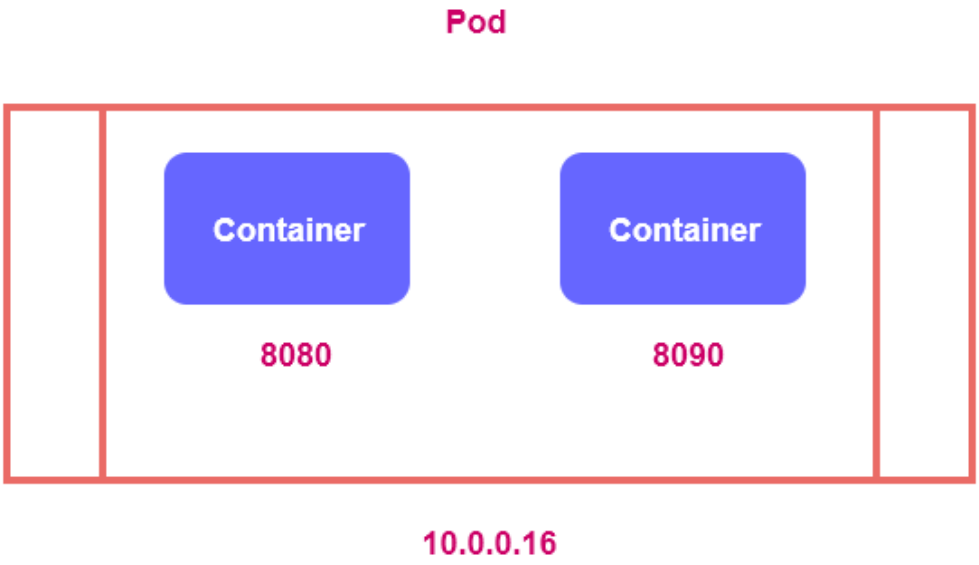

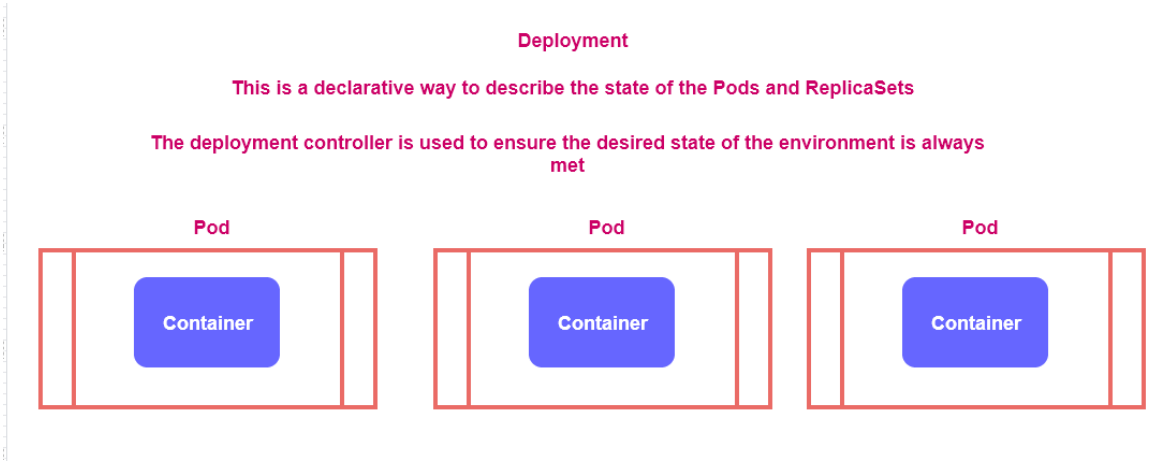

Develop for Azure Storage

Azure Storage Accounts

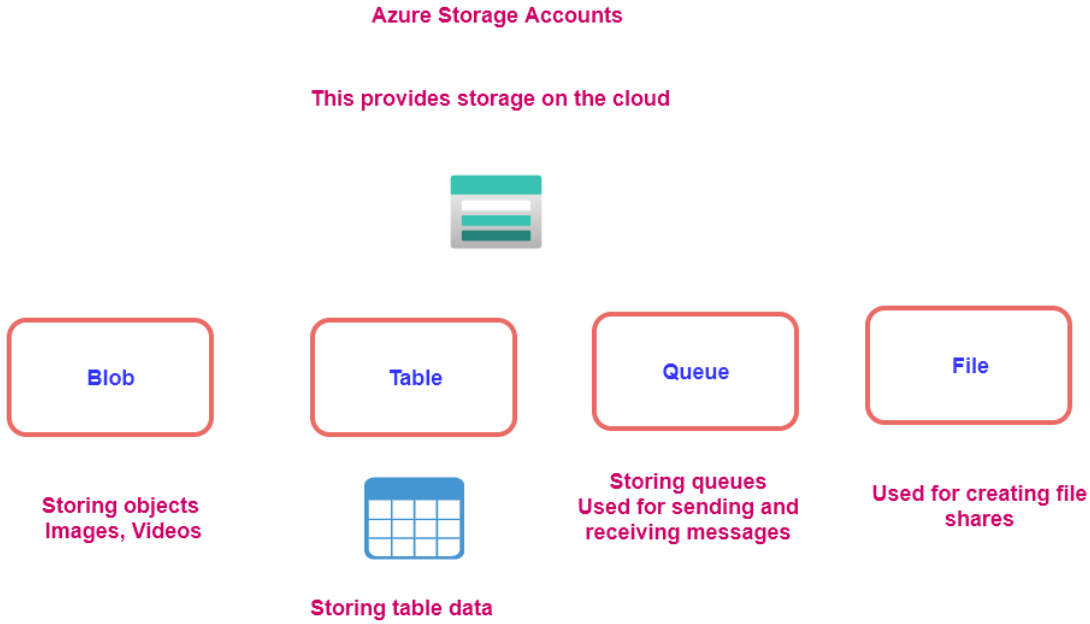

Azure Blob versioning

- This feature allows you to maintain previous versions of an object.
- Here you can restore to an earlier version of a blob to recover data if it is erroneously modified or deleted.
- For each subsequent modification of a blob, a newer version of the blob is created.
- Here the blob versions are immutable. They can't be modified.
- If you enable versioning, you can also disable it at any point in time. But this does not delete existing blobs, versions, or snapshots

Use lifecycle rules to transition objects from one tier to another

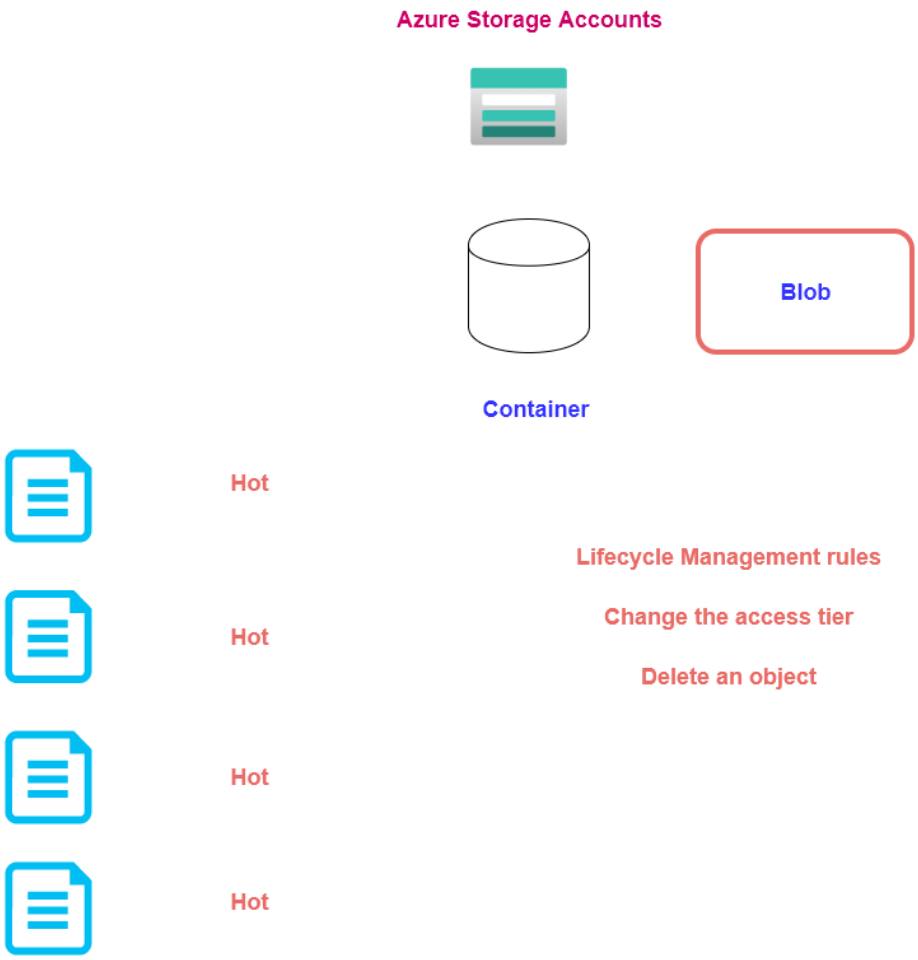

Azure Blob snapshots

- This takes a read-only version of a blob at a particular point in time.
- You can also make use of the blob versioning feature. This provides a better way to use a previous version of a blob.
- You can create as many snapshots as required of the blob.
- When you delete the blob, the snapshots also have to be deleted.

Azure Blob – Change Feed

- This feature provides transaction logs of all of the changes that occur to the blobs and the blob metadata in the storage account.
- Here you get an ordered, guaranteed, durable, immutable and read-only log of the changes.

- The changes events are recorded in the Apache Avro format specification.
- This file format provides a compact way to store data. And also, a fast way to access the data.
- Currently General Purpose v2 and Blob storage accounts support the change feed feature.

Get locks for your objects with the help of blob leases

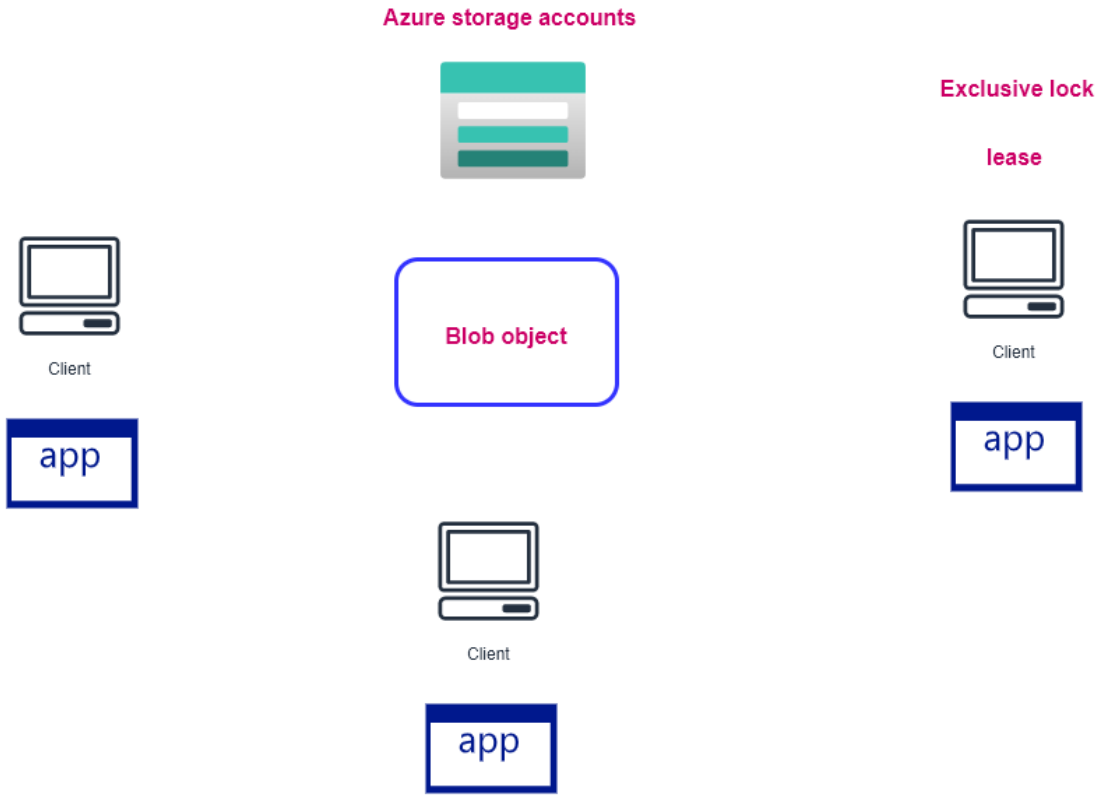

Azure Table storage

Azure Table storage

Table storage

Azure Storage Accounts

Data as tables

Stores non-relational structured data

NoSQL data

Here data is stored with the help of key/attributes and values

Follows a schemaless design

Azure Cosmos DB

Use Azure Cosmos DB as your NoSQL data store

Azure Cosmos DB

This is a fully managed NoSQL database

It provides fast response time

It provides high availability for your data

It is highly scalable

It carries out automatic updates and patching to the underlying database engine

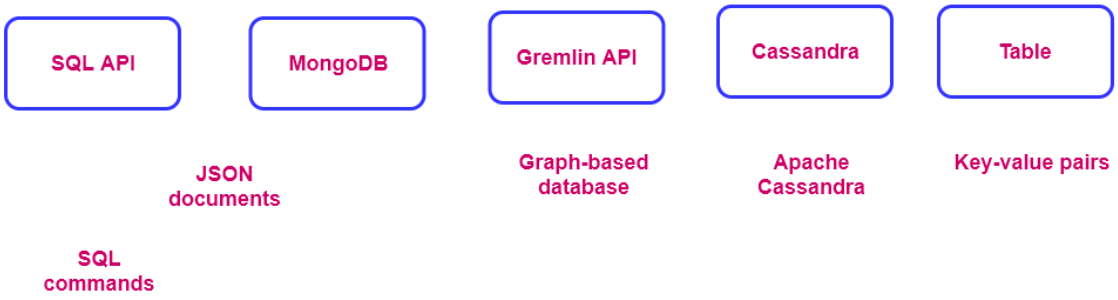

Consistency levels

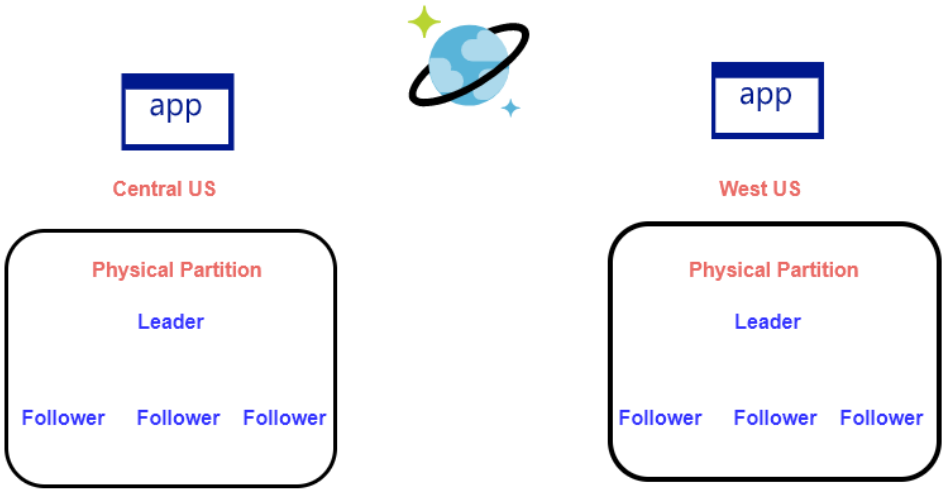


```
{
  "id" : 1,
  "name" : "John",
  "order-count" : "10"
}
```

```
{
  "id" : 1,
  "name" : "John",
  "order-count" : "12"
}
```

Consistency

1. Strong - You get consistency but loose on performance

2. Eventual - You win on performance but loose on consistency

3. Bounded staleness - Here the reads can lag behind the writes by at most "K" versions of an item or by "T" time interval

4. Session - Here reads can be guaranteed for the same session

5. Consistent Prefix - Here there is a delay in the reads of the most recent data but you will never see out of order writes

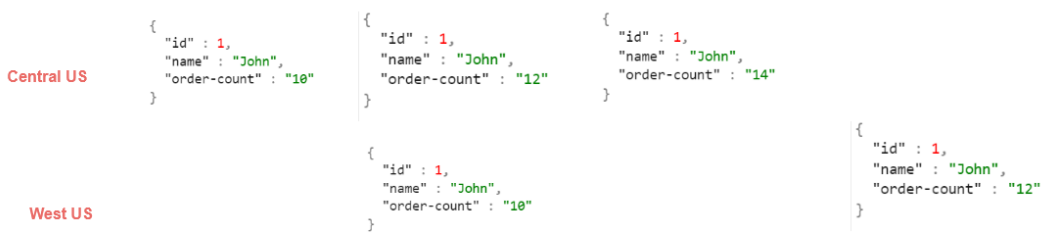

You can use synthetic keys for your partition key

Cosmos DB

Synthetic partition key

```
{
  "year" : 2020,
  "class" : "A"
}
```

```
{
  "year" : 2020,
  "class" : "A",
  "key" : "2018-A"
}
```

Combine keys

```
{
  "year" : 2020,
  "class" : "A",
  "key" : "2018-10"
}
```

Generate a random suffix

Embedding data

```
{
  "customerid": 2,
  "customername" : "Jane",
  "city" : "New York",
  "orders" :[
  {
  "orderid" : 1000,
  "productid" : 100,
  "quantity" : 2
  }]
}
```

1. Here we are denormalizing the data

When to embed data

1. When there are contained relationships between entities
2. When there is a one to few relationship between entities
3. When the embedded data changes infrequently
4. When the embedded data is queried frequently together

Referencing data

Seperate Document

```
{
  "customerid": 2,
  "customername" : "Jane",
  "city" : "New York",
  "orders" :[1000,1002]
}
```

Seperate Document

```
{
  "orderid" : 1000,
  "productid" : 100,
  "quantity" :2
}

{
  "orderid" : 1002,
  "productid" : 101,
  "quantity" :4
}
```

Referencing data

1. When representing one-to-many relationships

2. When representing many-to-many relationships

3. When there are related data changes

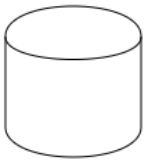

Container

Partition key

Change Feed, stored procedures and triggers

1. Change feed - This listens to all changes in a container and outputs a sorted list of documents that were changed in the order in which they were modified

Cosmos DB container

Change feed

Azure Function

Azure Blob storage

2. Stored Procedure - Here you can perform an operation at the database level

JavaScript

Stored Procedure

Cosmos DB container

Partitioning in Azure Cosmos DB

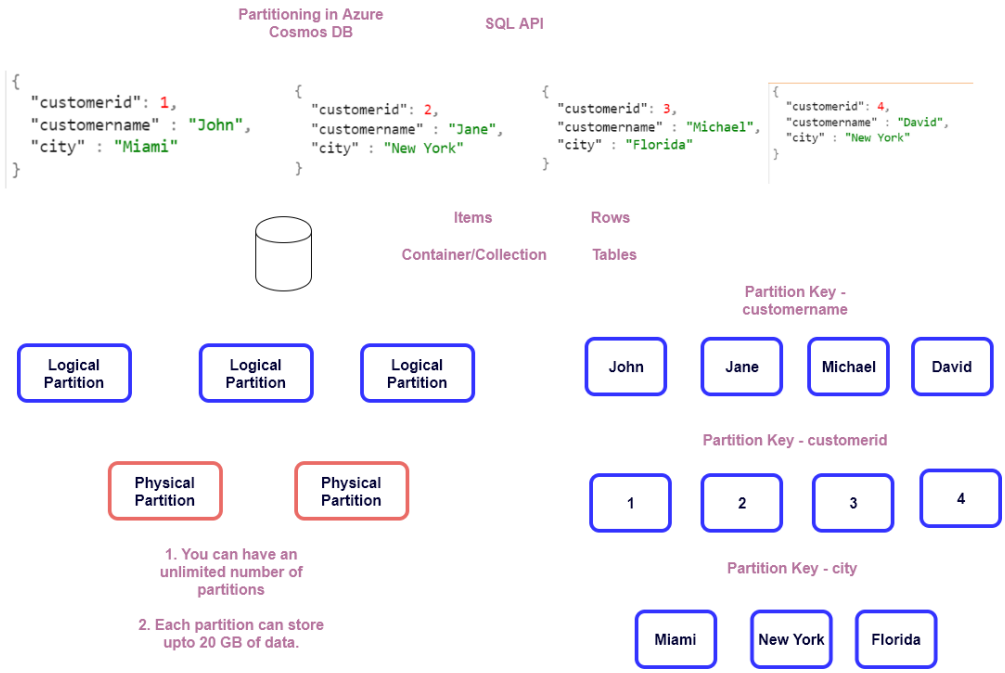

Time-to-live

- You can specify a time to live for all items in a container for just a specific item.
- After the time to live setting is reached, the item will be deleted automatically.
- The item will be deleted using the left-over request units.
- If there are no left-over request units, the deletion of the item will be delayed.

› You can set the value of the time to live at the container or item level.

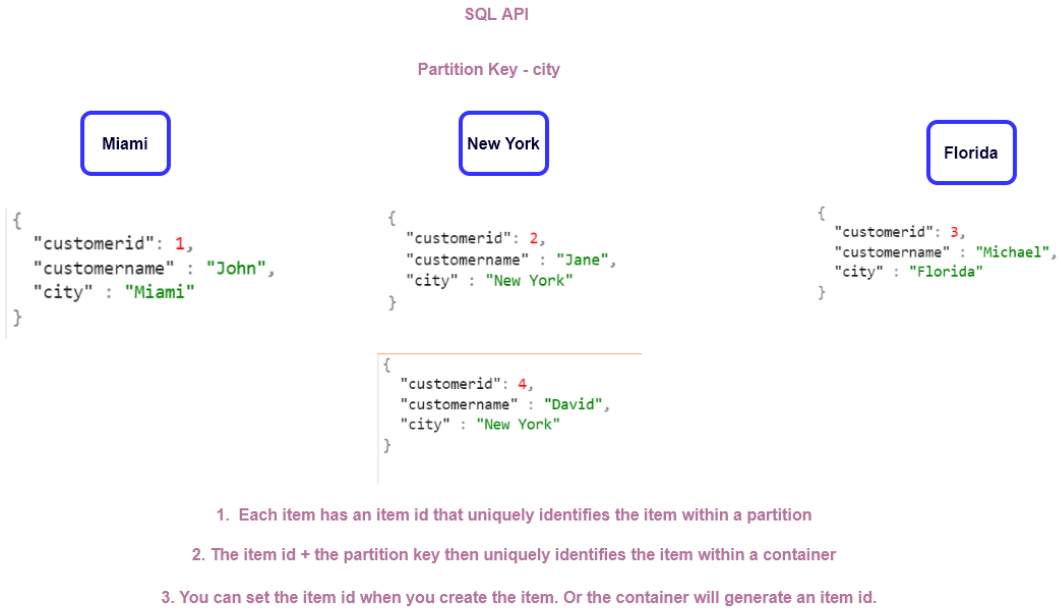

Implement Azure Security

Application Objects

Use Application Objects to represent your applications

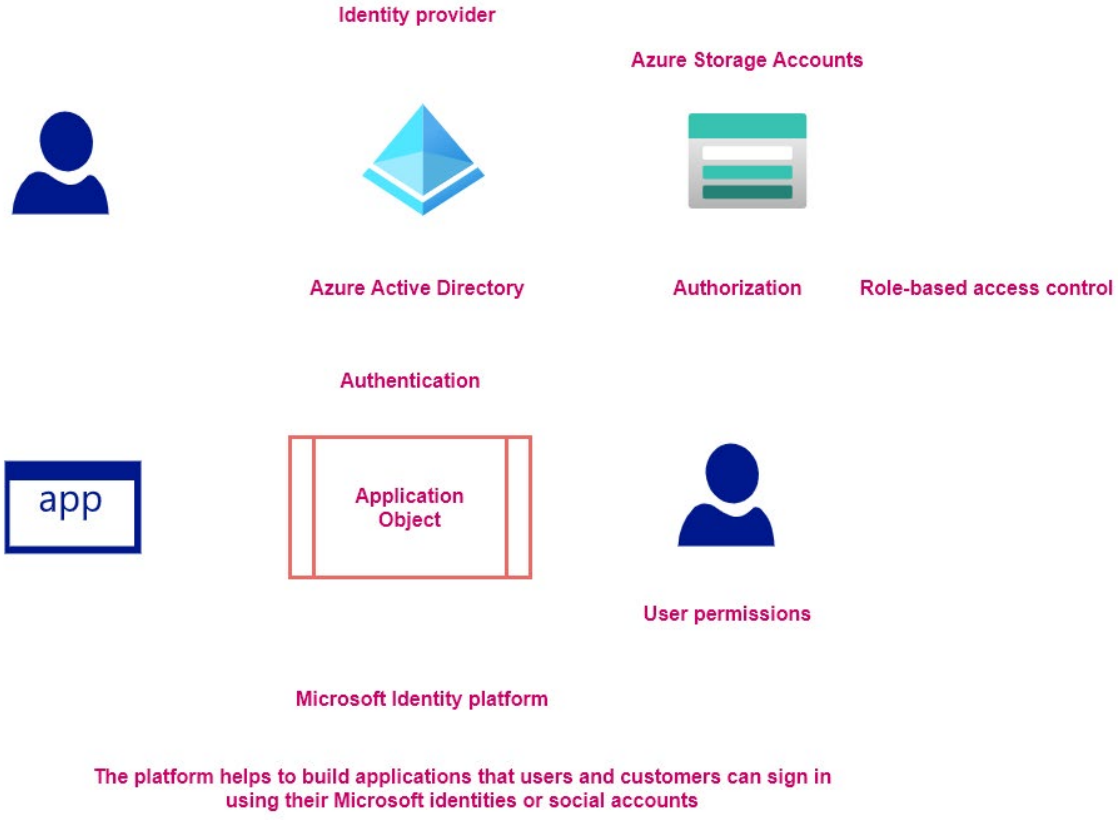

Using Managed identities

This helps to securely access resources in Azure

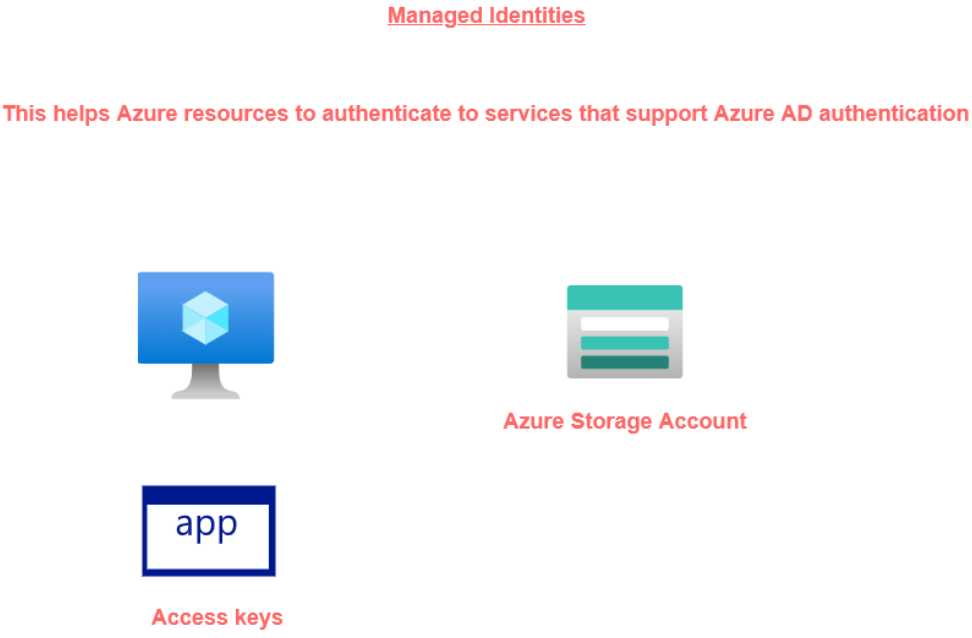

Reference link - https://docs.microsoft.com/en-us/dotnet/api/overview/azure/identity-readme

The basics of authentication

Using the POSTMAN tool to access blob storage

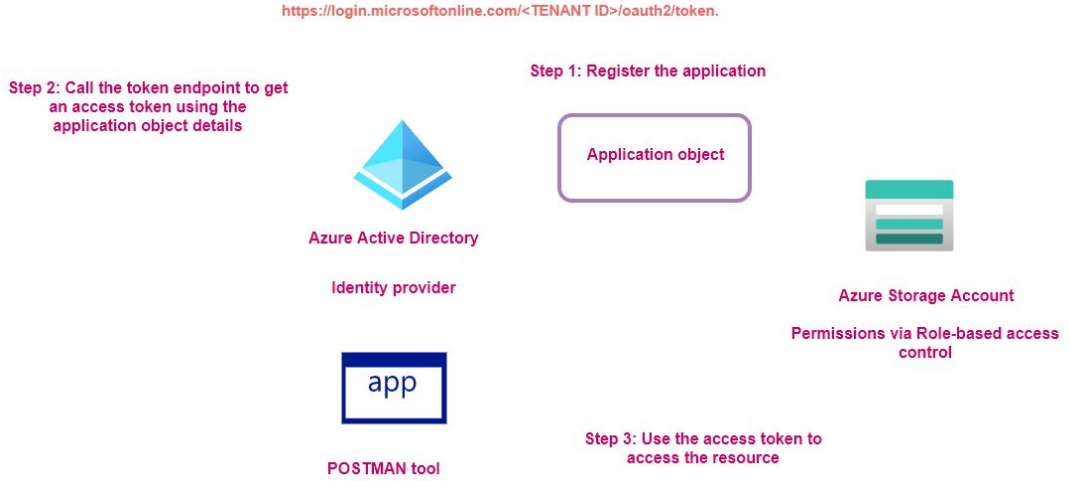

OAuth

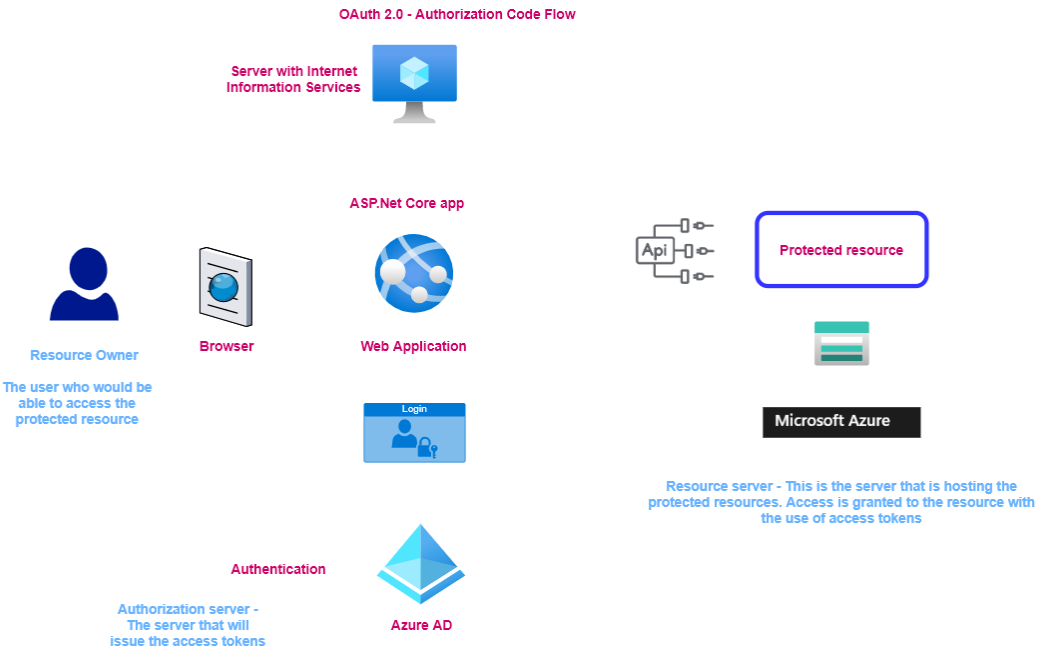

Getting the access token

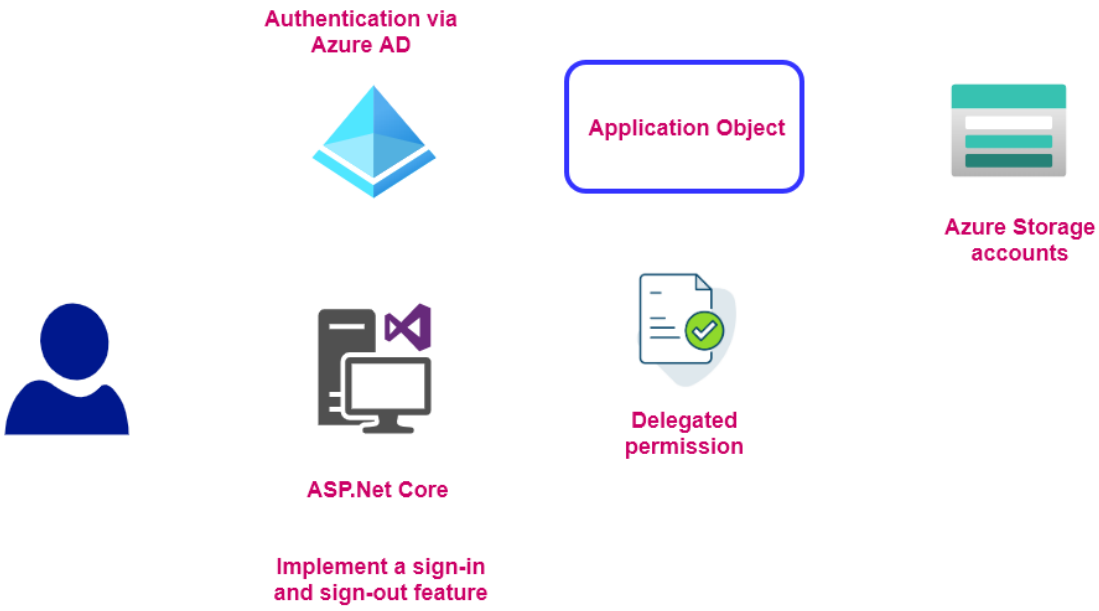

Azure AD Group claims

Azure AD

Role

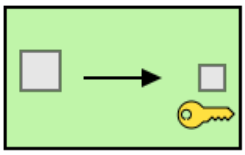

Check claims that are returned in the id token

Azure AD Groups

Protecting your API's

Authentication and authorization

Security

Ensure an access token is used to grant access to the API's

Protect your web api's

Security

Logging

Getting data

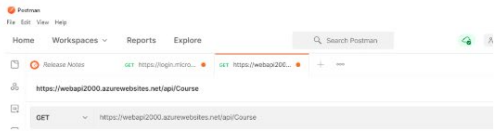

Application Object for the POSTMAN tool

Here we will be using the client credentials flow to get the access token

Application object for the web API

We will then create an App role

We will then ensure the POSTMAN aplication object has acces to the application object for the web api

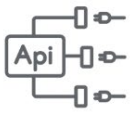

Course API

The application will have the required authorization modules

Azure AD

OAuth Implicit flow

Resource Owner

The user who would be able to access the protected resource

Browser

Client application - JavaScript , AngularJS

Web Application

Protected resource

Microsoft Azure

Resource server - This is the server that is hosting the protected resources. Access is granted to the resource with the use of access tokens

Authentication

Azure AD

Authorization server - The server that will issue the access tokens

Monitor, troubleshoot and optimize solutions

Application Insights

- Application Performance Management service for web developers.
- You can use this tool to monitor your applications.
- It can help developers detect anomalies in the application.

- It can help diagnose issues.
- It can also help understand how users use your application.
- It also helps you improve performance and usability of your application.
- **How does it work**
- You install a small instrumentation package within your application.
- You can see the statistics of your application locally in Visual Studio as you run your application.
- You can also use the Application Insights resource in Azure to monitor your application.

- **What are the different aspects monitored by Application Insights**
- Request rates, the response times and failure rates – This is done at the page level.
- Exception recorded by your application.
- Page views and their load performance as reported from the user's browser.
- User and session counts.
- Performance counters of the underlying Windows or Linux Machines.
- Diagnostic trace logs from your application.
- Any custom events or metrics that the developer writes themselves in the code.

- Understanding how your users use the application
- **Funnels** – You can create a funnel from one stage to another stage of your application.
- You can then see how users are progressing through the stages of the funnel.

- **User Flows** – This helps visualize how users navigate between pages in your site. This can help answer question such as
  - Does the user navigate away from a page on your site
  - What do users click on a page on your site
  - Where are the places where users churn most on your site
  - Are there places where users repeat the same action over and over

- **Impact** – This helps decide if a page is having an impact on your application.
- It can help answer the question as to whether the page load time is impacting how many people convert on a page in the application.
- **Retention** – This helps you understand how many users return to your application.
- It can also help understand if users are able to perform certain tasks in your application.

Azure Front Door Caching and Compression

By default, Azure Front Door has the capability to deliver large files. If a file is large, it splits the file into chunks of 8 MB. If the file is not in the cache, the large file is retrieved from the origin. Once the first 8 MB chunk is received , it is sent to the user and then in parallel , the next 8 MB chunk is retrieved from the origin.

Azure Front Door also supports compression. It supports all of the following MIME types. But the file needs to be in size between 1 KB and 8 MB. It supports both gzip and Brotli compression.

- "application/eot"
- "application/font"
- "application/font-sfnt"
- "application/javascript"
- "application/json"

- "application/opentype"
- "application/otf"
- "application/pkcs7-mime"
- "application/truetype"
- "application/ttf",
- "application/vnd.ms-fontobject"
- "application/xhtml+xml"
- "application/xml"
- "application/xml+rss"
- "application/x-font-opentype"
- "application/x-font-truetype"
- "application/x-font-ttf"
- "application/x-httpd-cgi"
- "application/x-mpegurl"
- "application/x-opentype"
- "application/x-otf"
- "application/x-perl"
- "application/x-ttf"
- "application/x-javascript"
- "font/eot"
- "font/ttf"
- "font/otf"
- "font/opentype"
- "image/svg+xml"
- "text/css"
- "text/csv"
- "text/html"
- "text/javascript"
- "text/js", "text/plain"
- "text/richtext"
- "text/tab-separated-values"
- "text/xml"
- "text/x-script"
- "text/x-component"
- "text/x-java-source"

Connect to and consume Azure and third-party services

Queues

The entire purpose of queues is to decouple components of your application

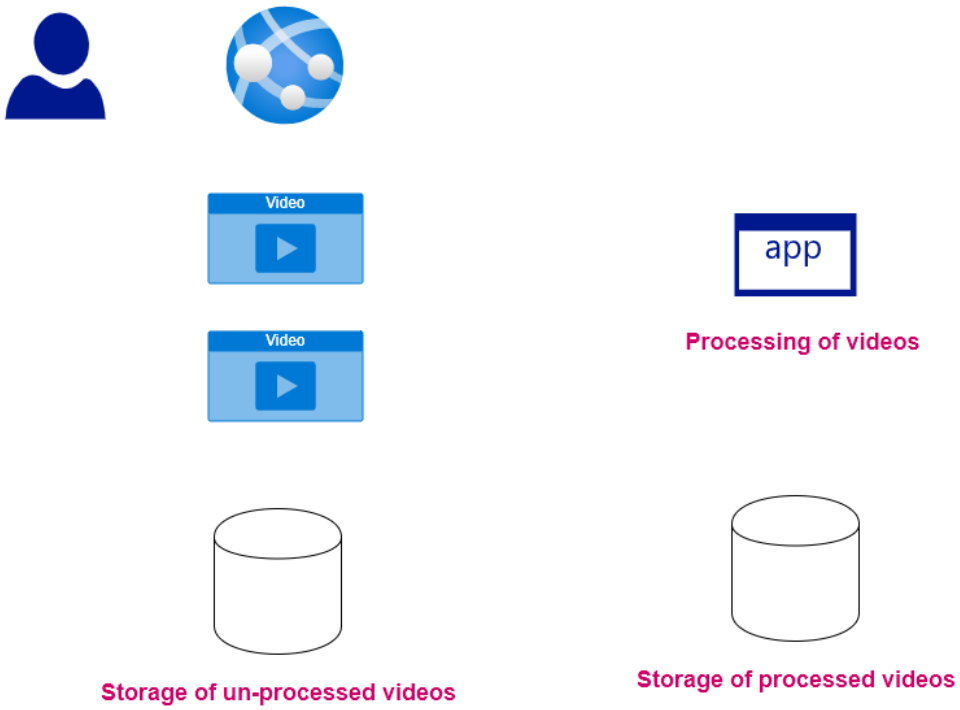

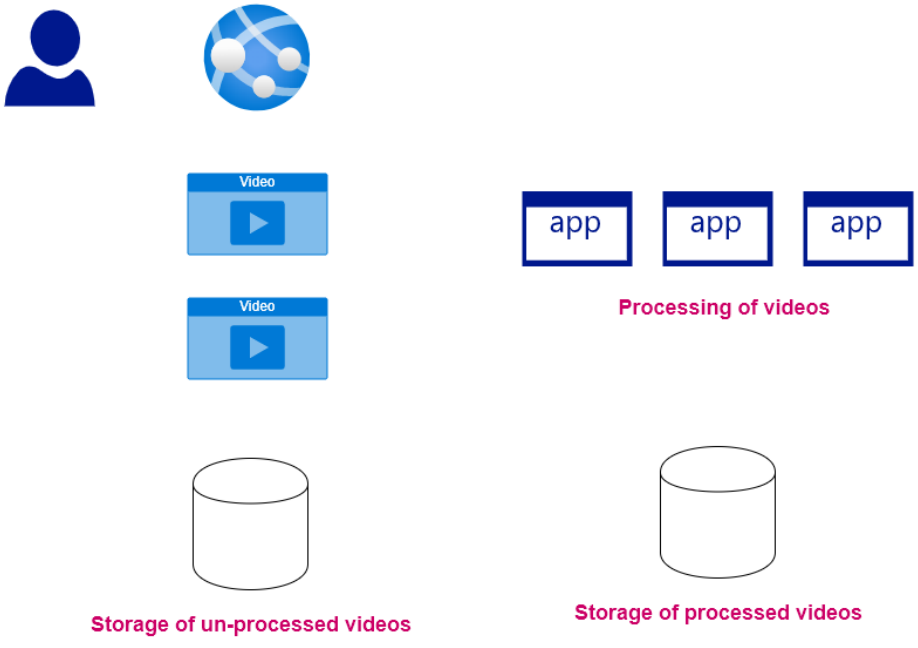

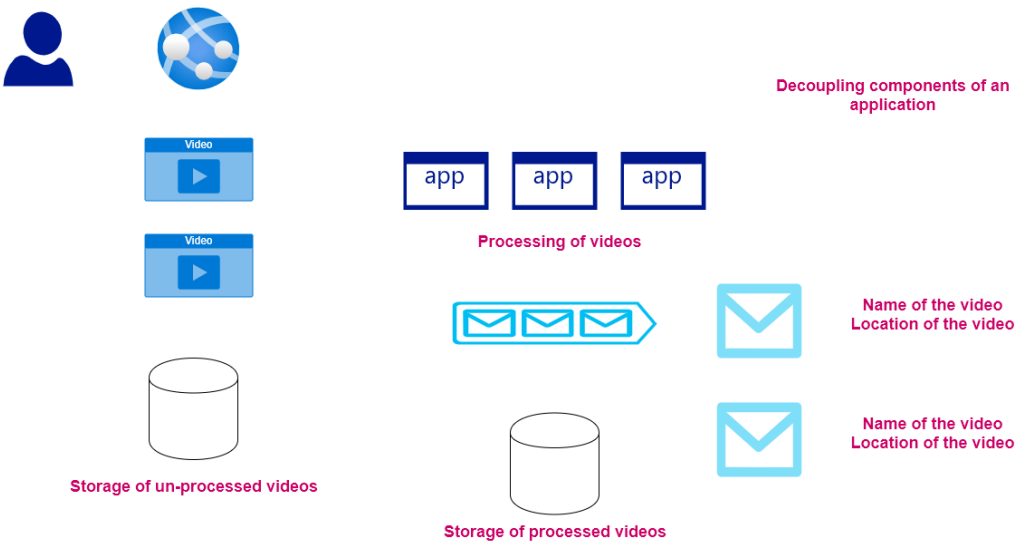

Azure Service Bus

In the PeekLock mode, you can mention what is the duration to lock messages

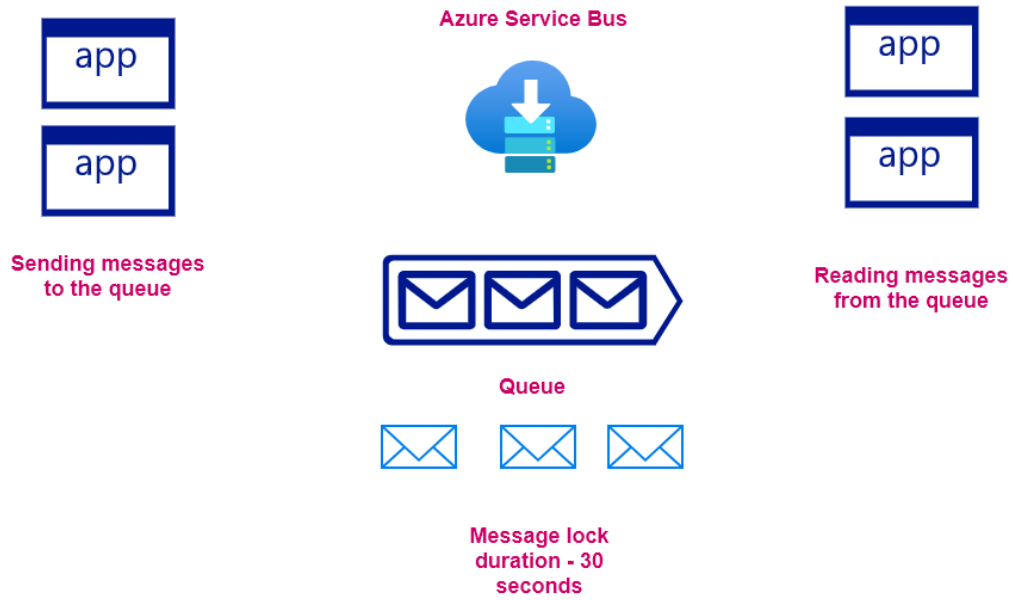

Messages that can’t be processed , can be taken from the dead-letter queue

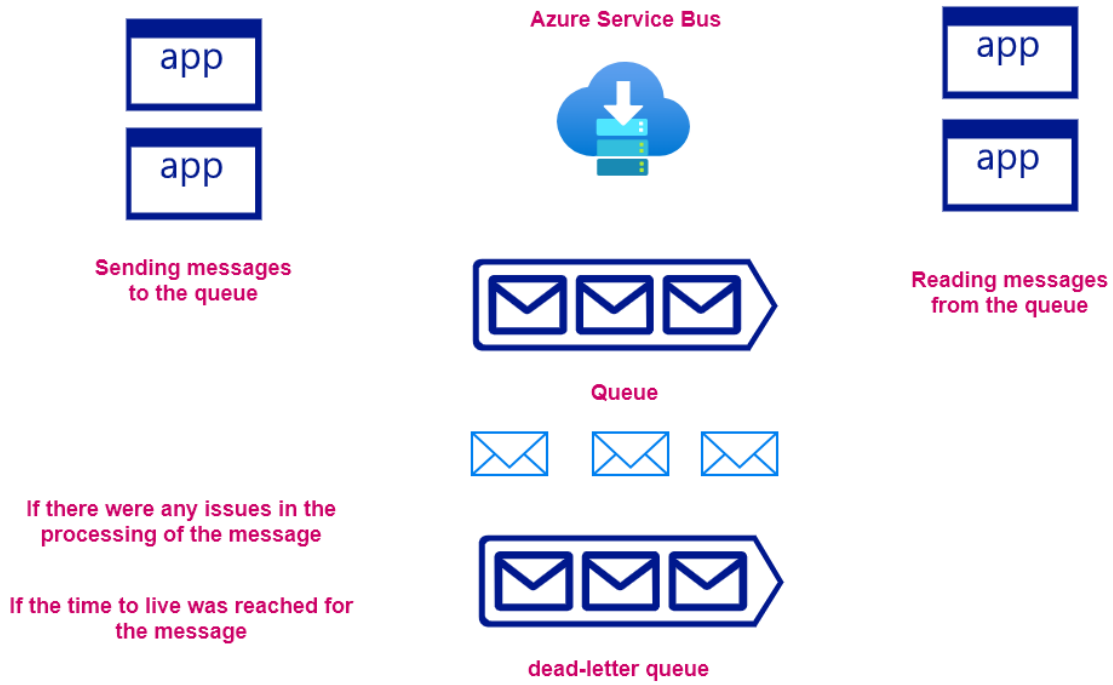

You can also detect duplicate messages from Azure Service Bus queues

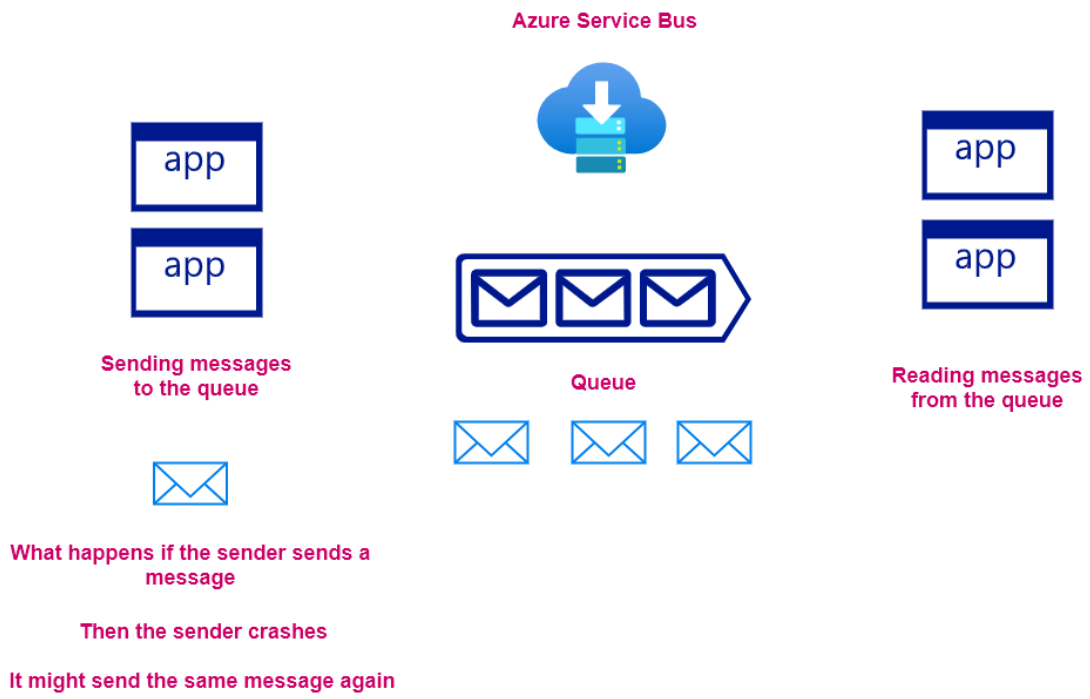

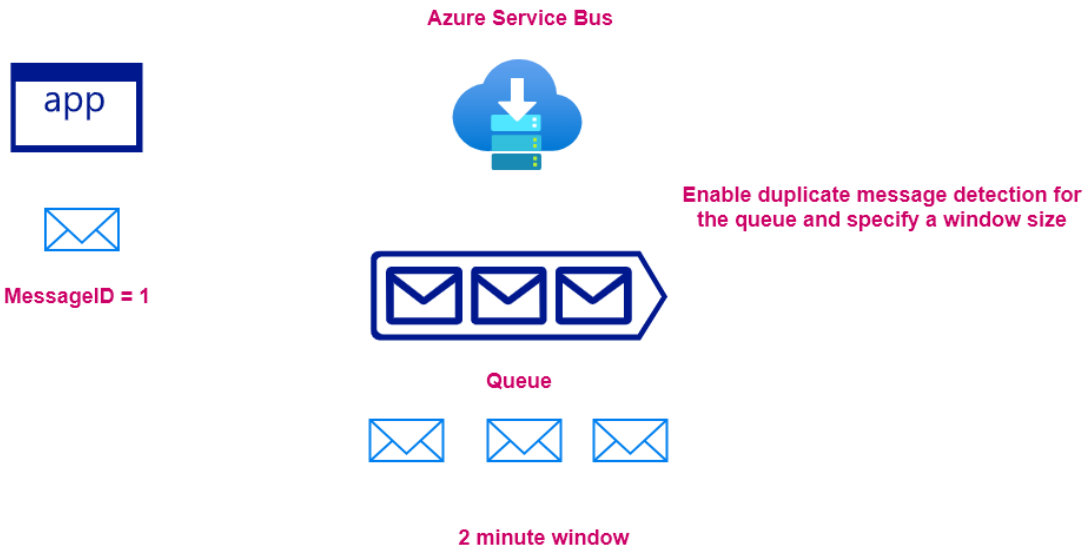

Understanding the different parts of a message

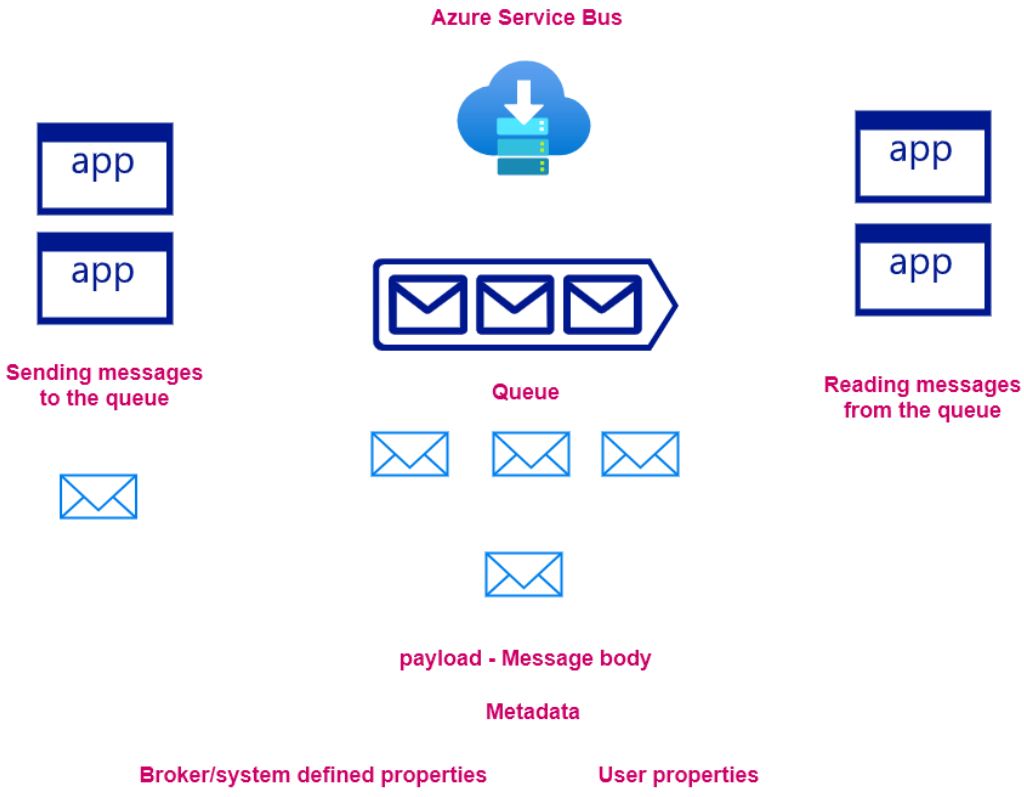

Azure Service Bus topics

Here subscribers can subscribe to get messages from the topic

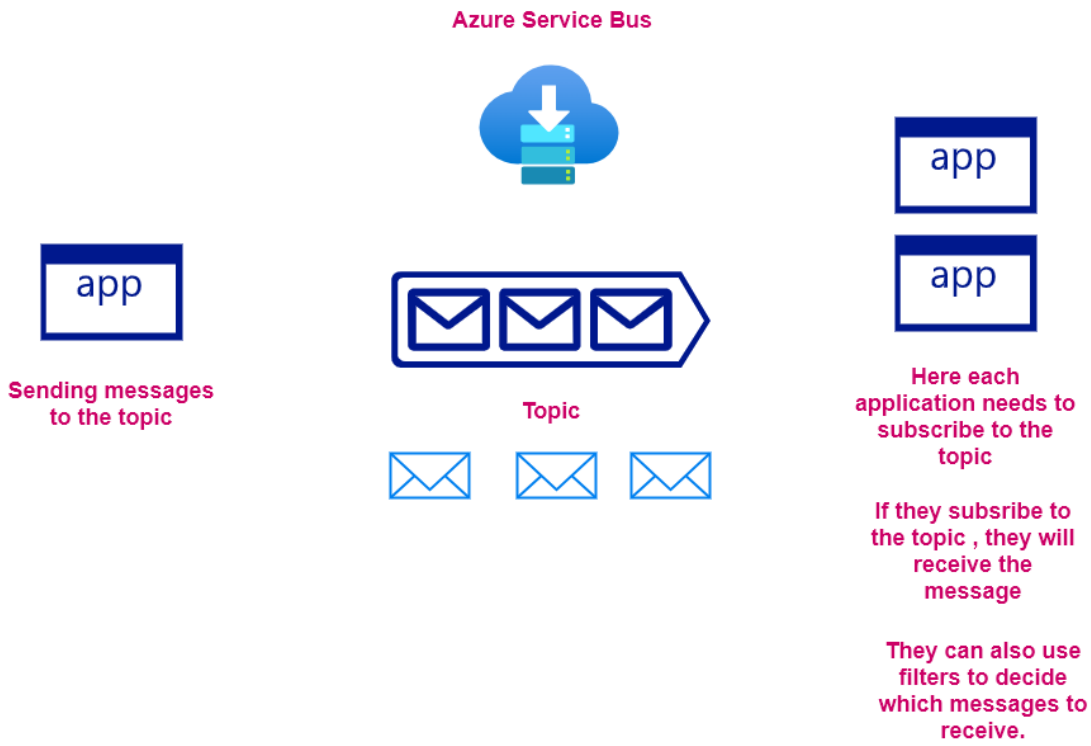

Azure Service Bus Topic filters

- Subscriptions can use filters to decide on which messages they want to receive.
- SQL Filters – Here you can use SQL-like conditions for the evaluation. This can be done against the system or user defined properties.

- For system-defined properties , the properties must be prefixed with sys.
- Boolean filters – This is a TrueFilter or FalseFilter.
- Correlation Filters – Here conditions can be used to match against the user's or system properties.
- Note:- Correlation Filters process faster than SQL filters.

Difference when it comes to Azure storage queues and Service Bus queues

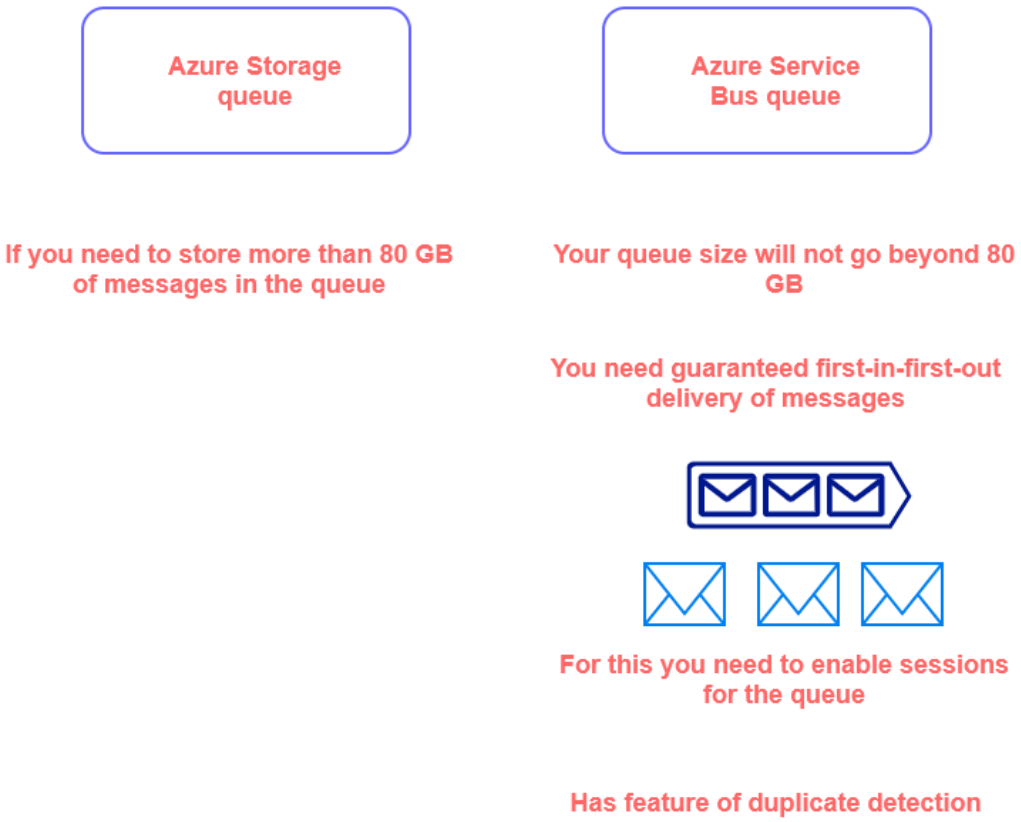

Azure Event Hubs

When you want to ingest large amounts of data , look towards using Azure Event Hubs

You can stream log data , telemetry data, any sort of events to Azure Event Hubs

Event Hubs Architecture

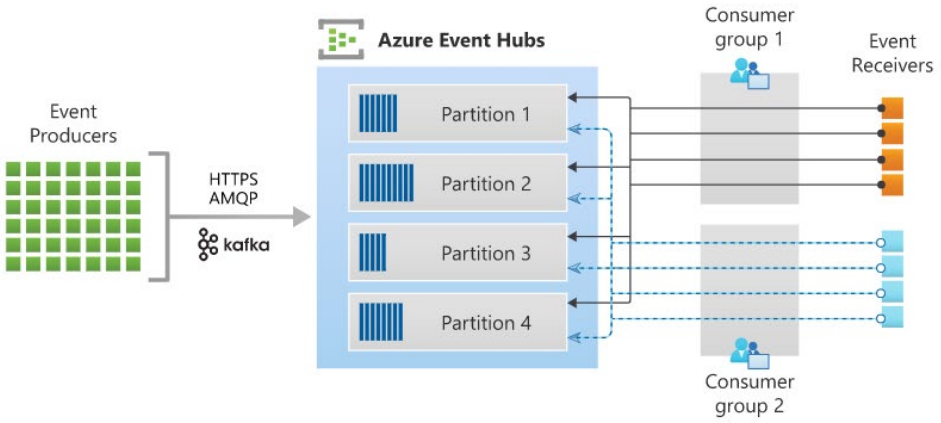

The different components

Event producers - This is an entity that sends data to an event hub. The events can be published using the protocols - HTTPS, AMQP, Apache Kafka

Partitions - The data is split across partitions. This allows for better throughput of your data onto Azure Event Hubs

Consumer groups - This is a view (state,position or offset) of an entire event hub

Throughput - This controls the throughput capacity of Event Hubs

Event Receivers - This is an entity that reads event data

When you send events or data onto Azure Event Hub, they are sent to a partition. The receiver can then receive events via a consumer group

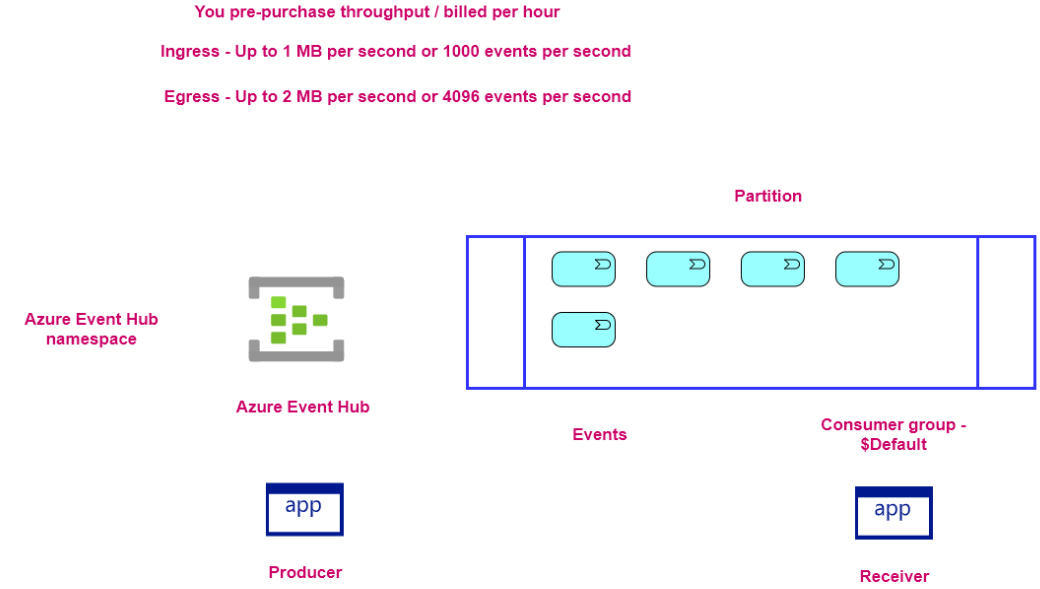

You can use the EventProcessor class to track the events which have been consumed from Azure Event Hub

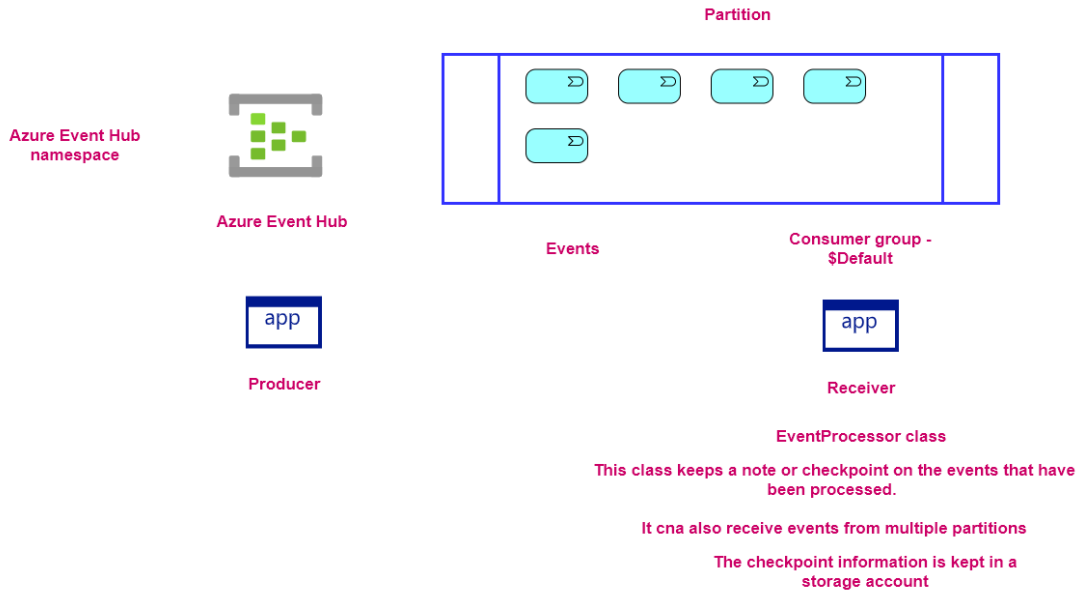

Recommendations when it comes to readers and partitions

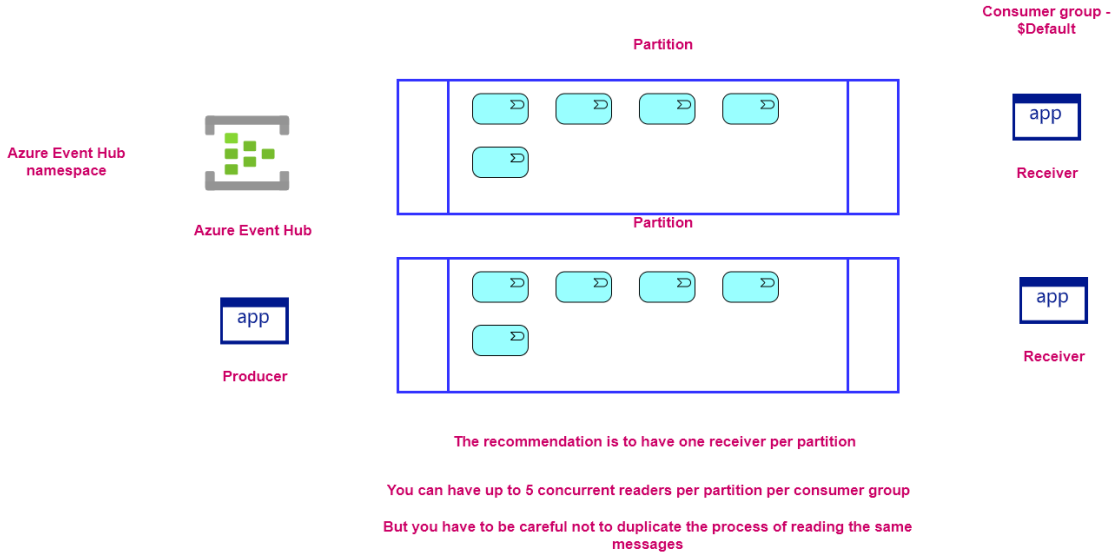

Azure Event Grid

Capture events from various Azure resources with the use of Azure Event Grid

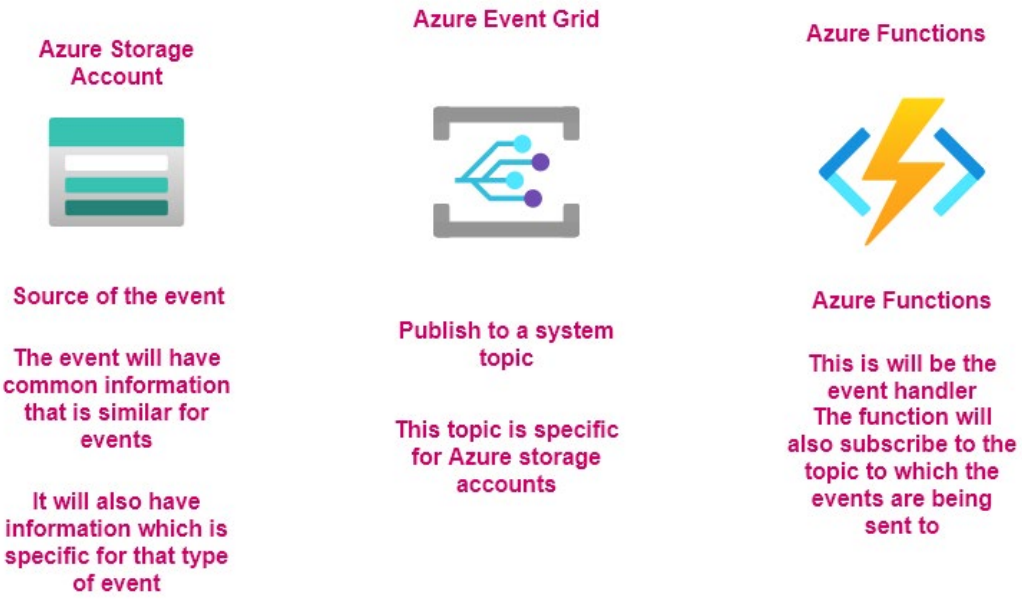

Use Azure Event Grid to capture events from an Azure storage account and then use Azure functions as the endpoint to capture the events.

To debug your Azure Event Grid code from Visual Studio in Azure Functions, make sure to use the ngrok tool

https://ngrok.com/

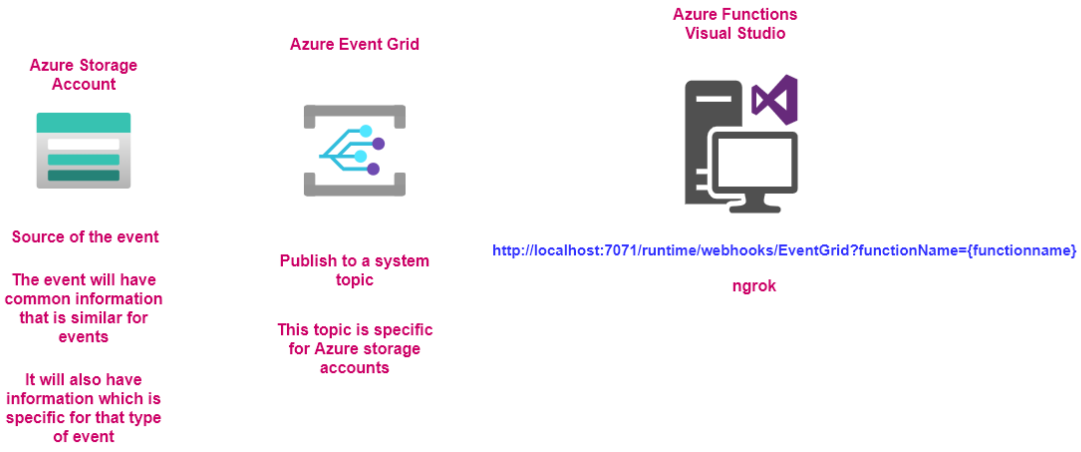

When working with Azure Event Grid, its important to know the schema. Here you know the format of the data that is sent from Azure Event Grid

Reference - https://docs.microsoft.com/en-us/azure/event-grid/event-schema

JSON

```
[
  {
    "topic": string,
    "subject": string,
    "id": string,
    "eventType": string,
    "eventTime": string,
    "data":{
      object-unique-to-each-publisher
    },
    "dataVersion": string,
    "metadataVersion": string
  }
]
```

The event sources send the events to the Azure Event Grid service in an array

The array can have several event objects

The total size of the array will be up to 1 MB

Remember that if you have your own custom Webhook, then you need to have the handshake code in place to take the validationcode and send it back to the validationURL.

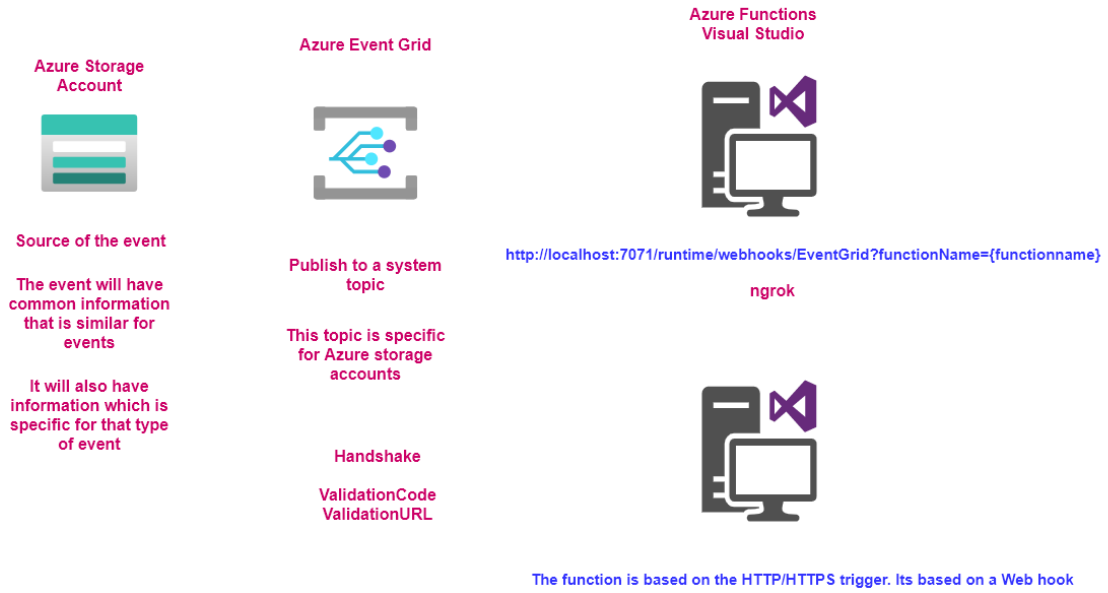

You can also create your custom topics

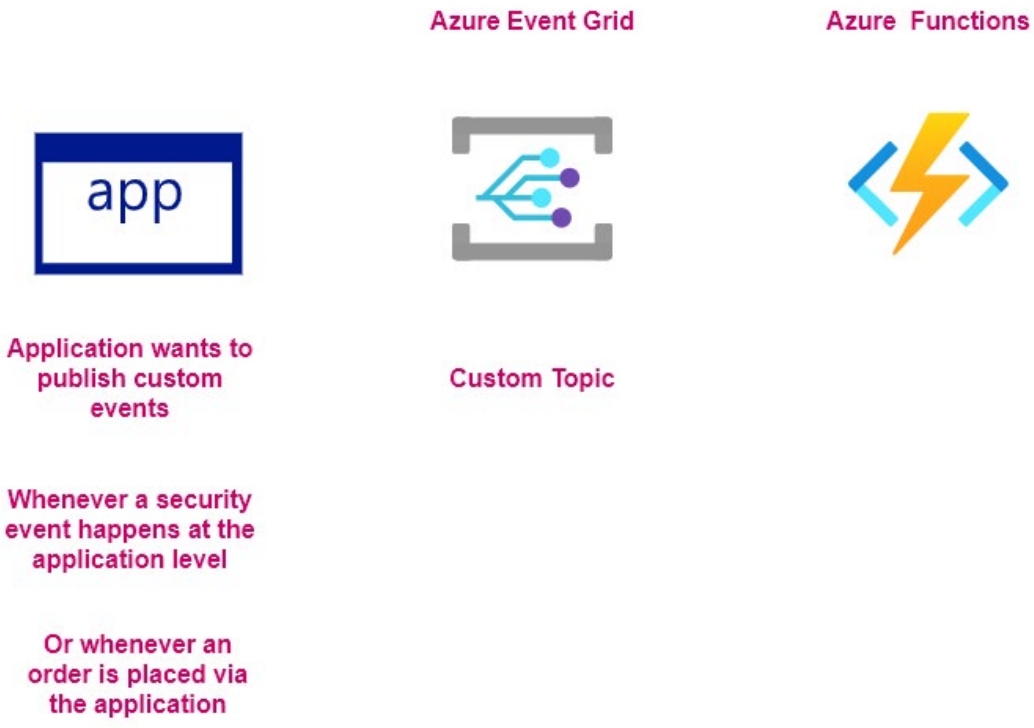

Comparison between the different messaging services

Azure Event Hubs

Used for telemetry or distributed data streaming

Message size - 256KB for Basic tier and 1 MB for Standard tier

Azure Event Grid

Events - Lightweight notification of a condition or a state change

Maximum size of an event is 1 MB

Azure Service Bus

Message - This is raw data that needs to be consumed.

Examples of use cases - Order processing and financial transactions

Message size - 256KB or 1 MB if you choose Premium messaging

Devices sending telemetry data

Azure Event Hubs

Process the data

When a change occurs to a blob

Azure Event Grid

Respond to an event notification

endpoint

Some differences when it comes to Storage Blob triggers and Azure Event Grid

Azure storage accounts
Blob service

When a new or updated blob is detected

Blob storage trigger

Azure functions

Azure Event Grid

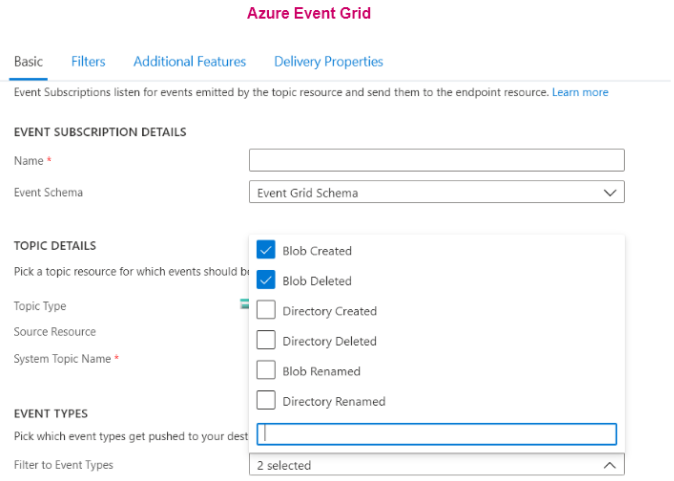

Azure Event Grid also has support for Blob-only storage accounts

Also if you need high-scale, consider using Event Grid over the Blob trigger - This is if you have more than 100,000 blobs in the storage account or if there are 100 blob updates per second.

Azure Logic Apps

Helps you to build workflows in Azure

Azure Logic Apps

WORKFLOW

Azure Functions

Developers define logic to perform some steps

When an administrative action is performed on a virtual machine

Email an administrator